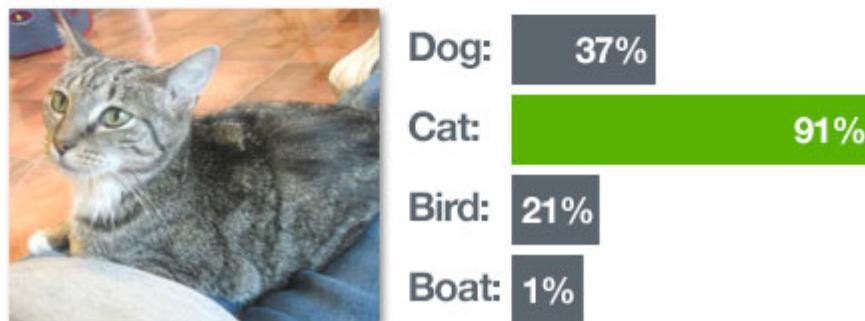
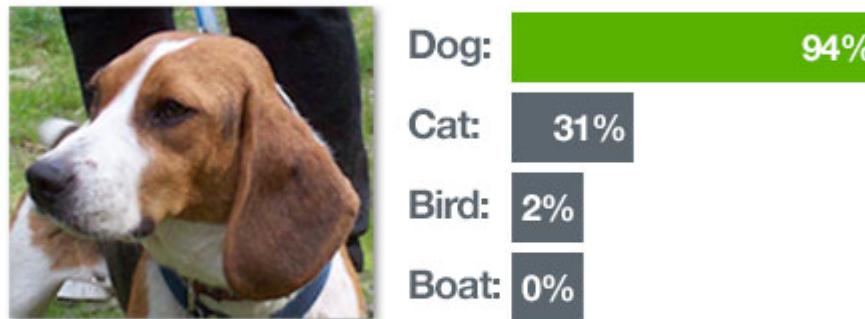
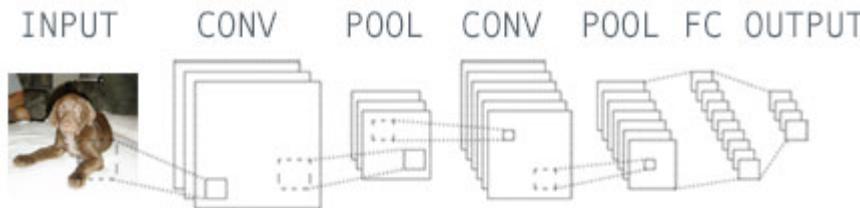


# Convolutional Networks



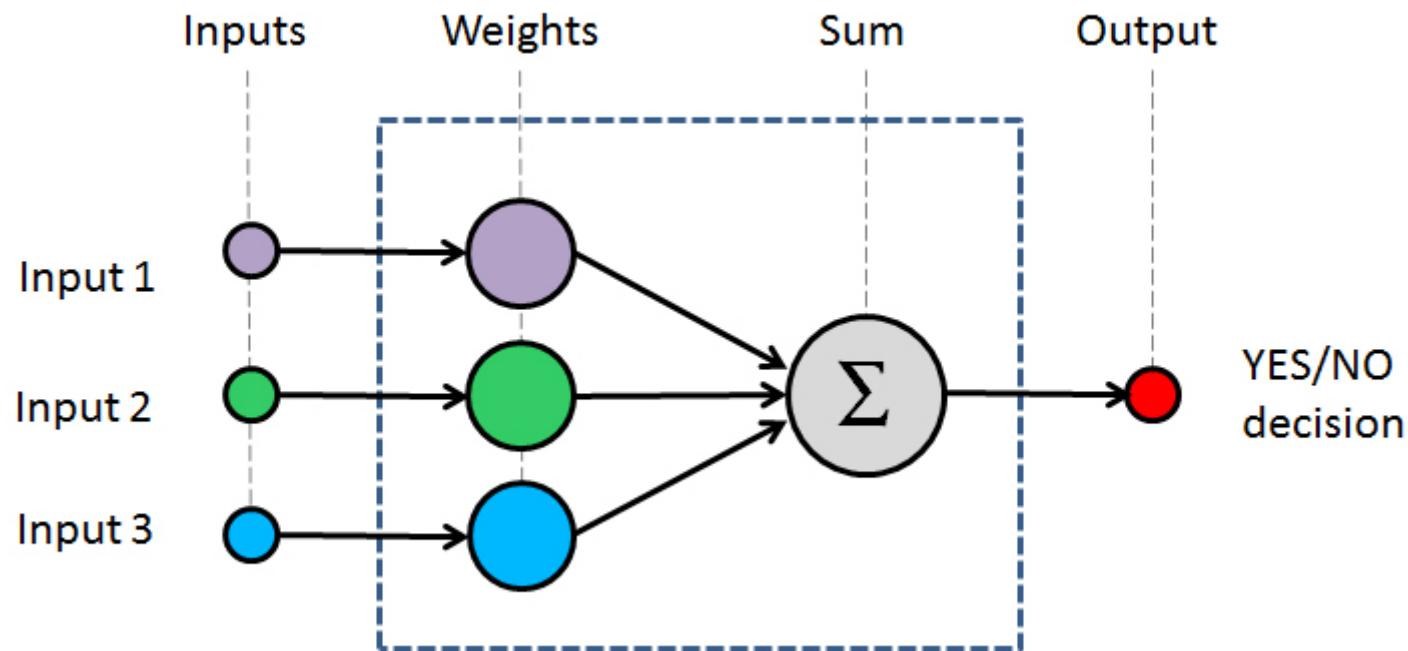
*The Truth About Cats & Dogs*

*Tony Reina*

# Goals

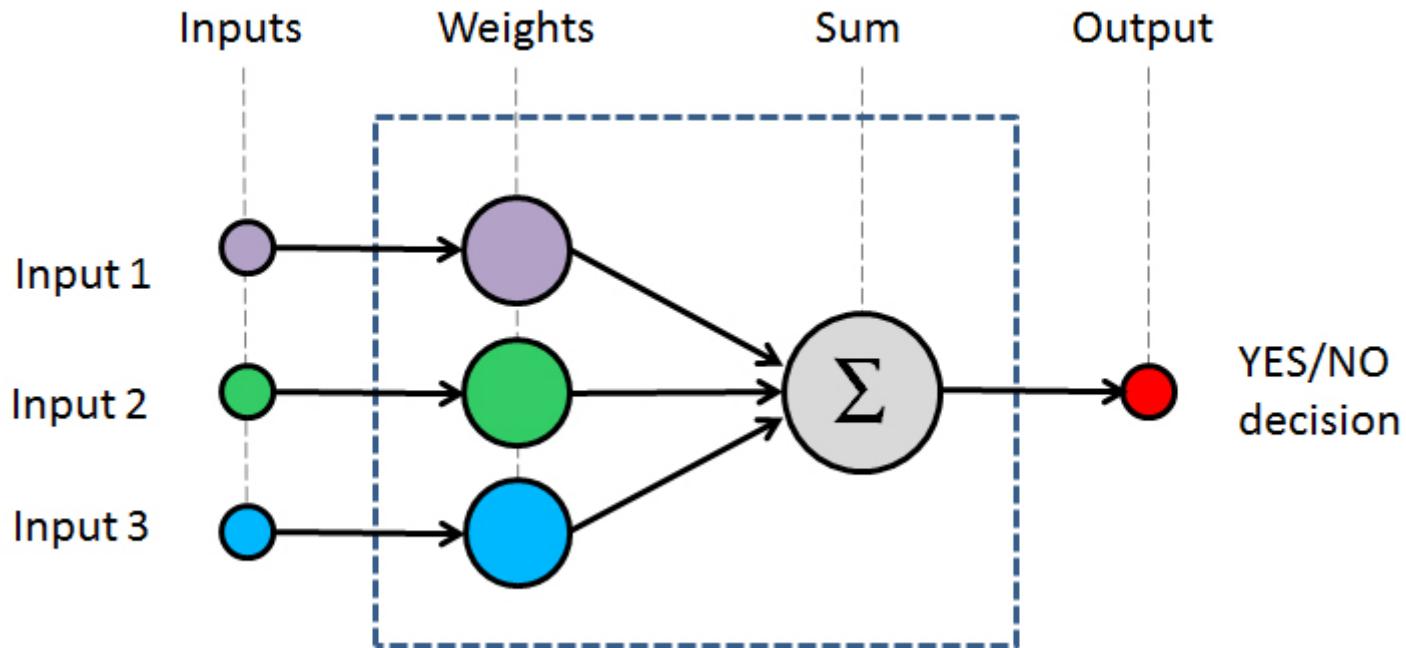
- What is convolution?
- Why should I care?
- How do I use it?

# What is convolution?



$$X \bullet W = \sum_{i=1}^n x_i \times w_i$$

# What is convolution?

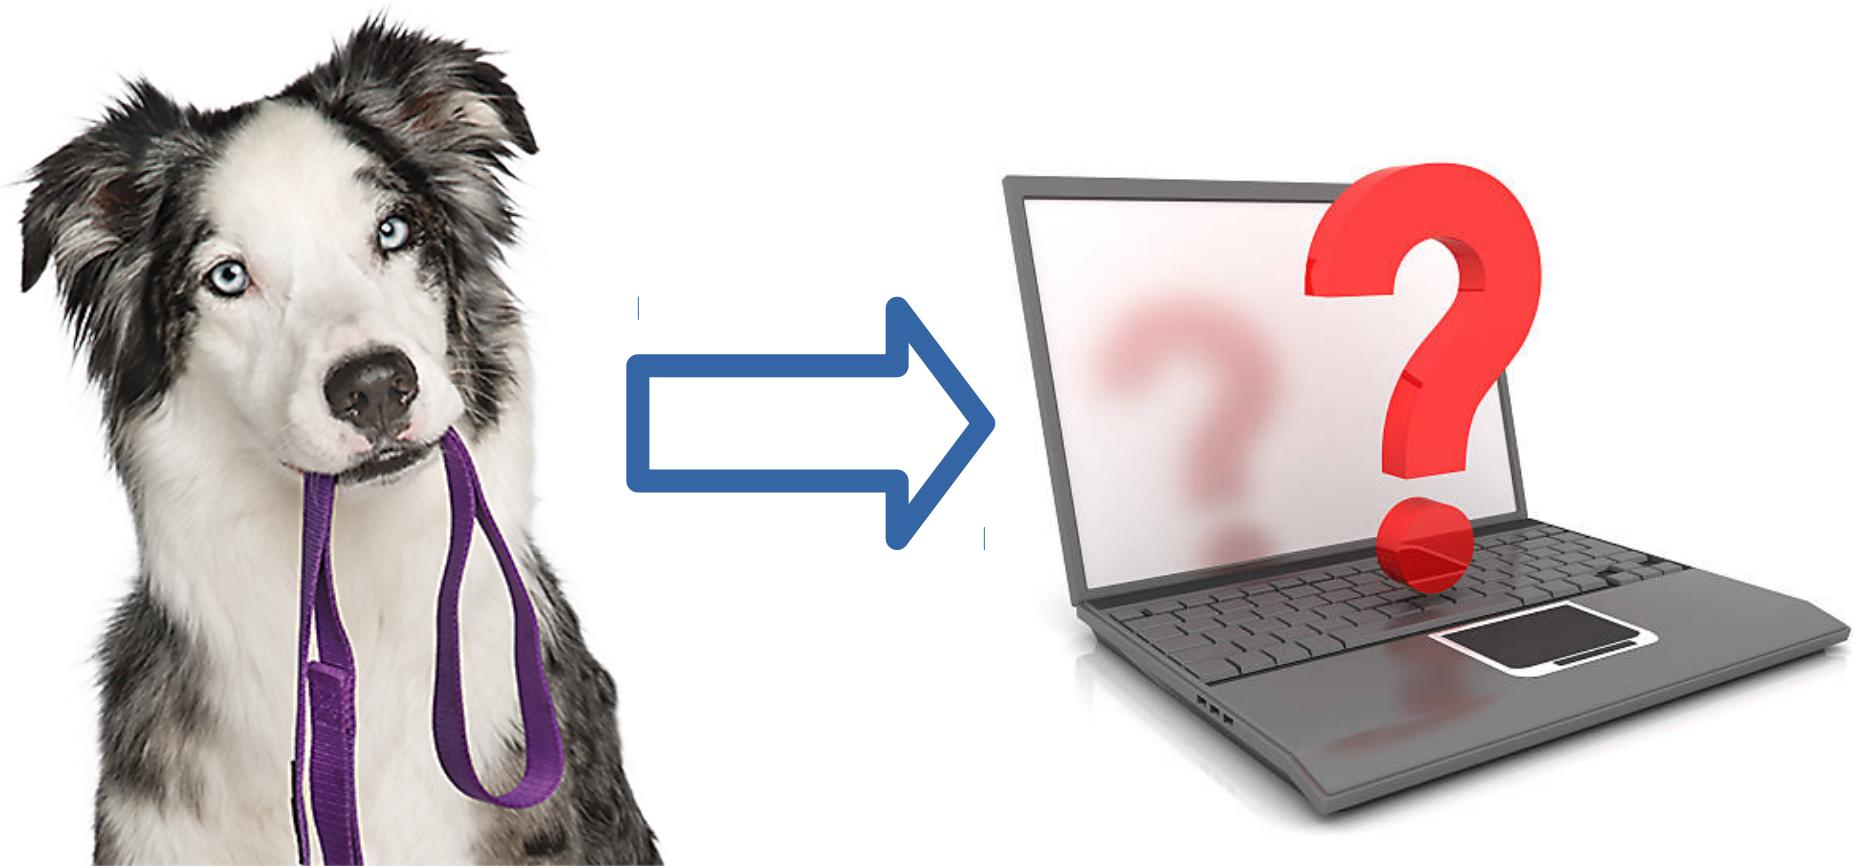


Out = 0

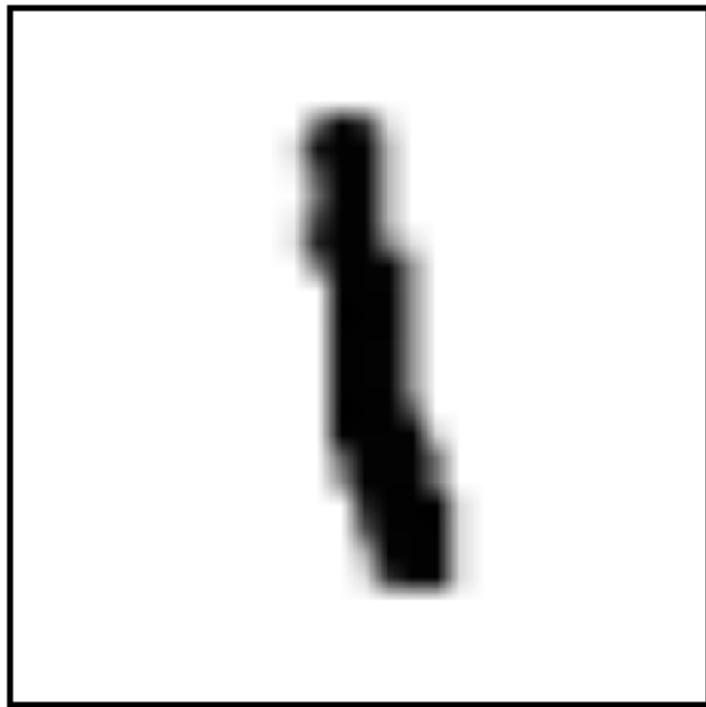
for index in number\_of\_inputs:

    Out = Out + (input[index] x weight[index])

# What is convolution?

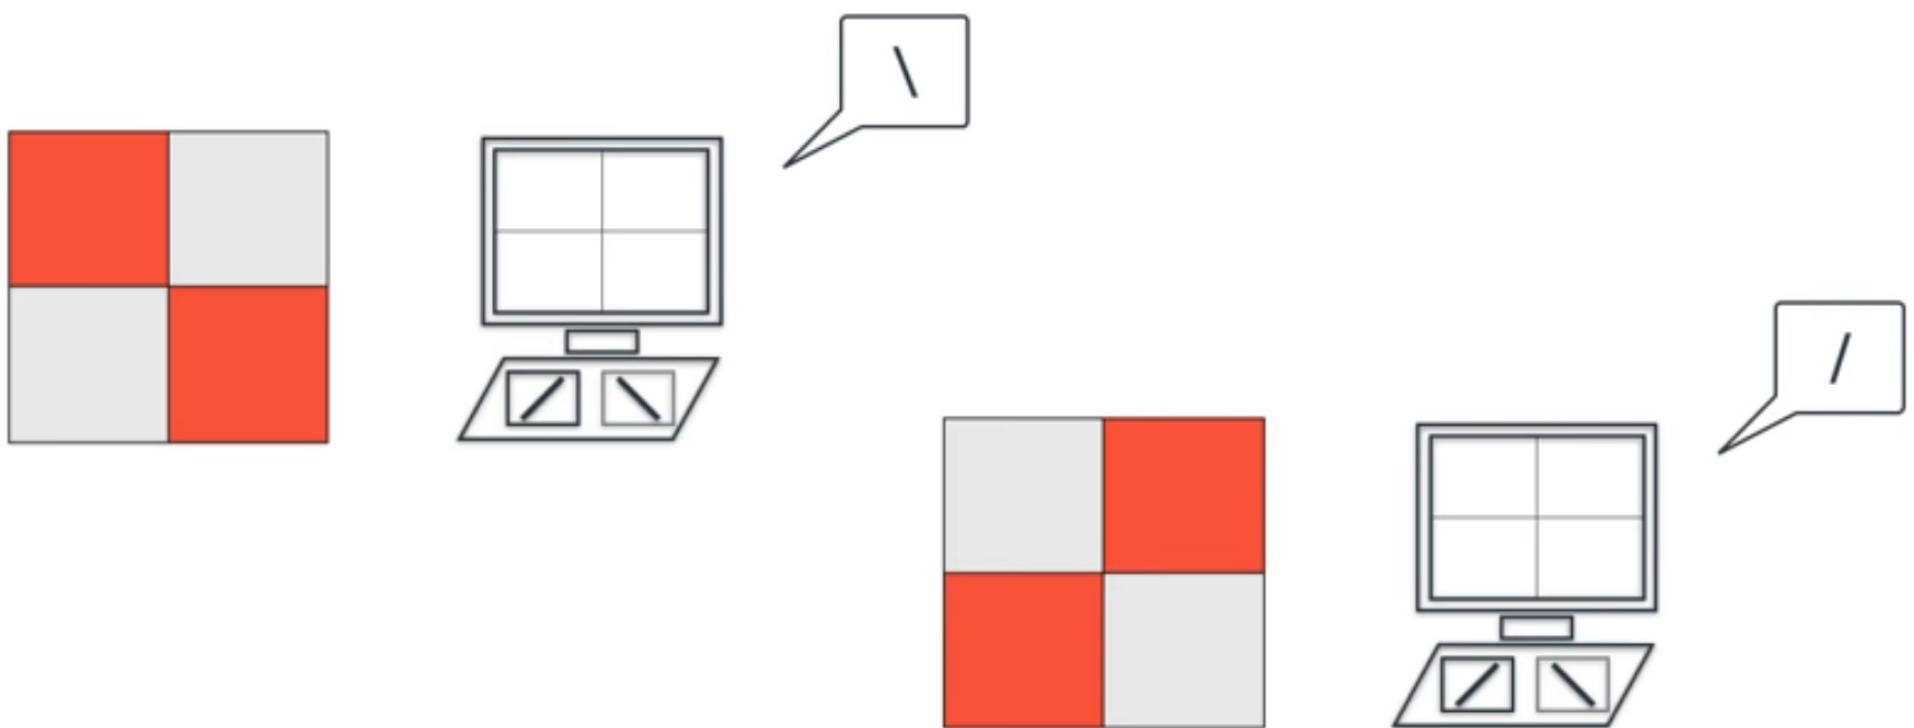


# What is convolution?



2

# What is convolution?



SOURCE: Luis Serrano, Udacity, <https://www.youtube.com/watch?v=2-Ol7ZB0MmU>

# What is convolution?

1	-1
-1	1

1	-1	-1	1
---	----	----	---

-1	1
1	-1

-1	1	1	-1
----	---	---	----

SOURCE: Luis Serrano, Udacity, <https://www.youtube.com/watch?v=2-Ol7ZB0MmU>

# What is convolution?

$$\backslash \begin{bmatrix} 1 & -1 \\ -1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & + & -1 & + & -1 & + & 1 \end{bmatrix} = 0$$

$$/ \begin{bmatrix} -1 & 1 \\ 1 & -1 \end{bmatrix}$$

$$\begin{bmatrix} -1 & + & 1 & + & 1 & + & -1 \end{bmatrix} = 0$$

SOURCE: Luis Serrano, Udacity, <https://www.youtube.com/watch?v=2-Ol7ZB0MmU>

# What is convolution?

$$\backslash \begin{bmatrix} + & - \\ 1 & -1 \\ - & -1 \\ -1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} + & - & - & + \\ 1 & -1 & -1 & 1 \\ +1 & +1 & +1 & +1 \end{bmatrix} = 4$$

$$/ \begin{bmatrix} + & - \\ -1 & 1 \\ - & -1 \\ 1 & -1 \end{bmatrix}$$

$$\begin{bmatrix} + & - & - & + \\ -1 & 1 & 1 & -1 \\ -1 & -1 & -1 & -1 \end{bmatrix} = -4$$

SOURCE: Luis Serrano, Udacity, <https://www.youtube.com/watch?v=2-Ol7ZB0MmU>

# What is convolution?

Convolutional Kernel

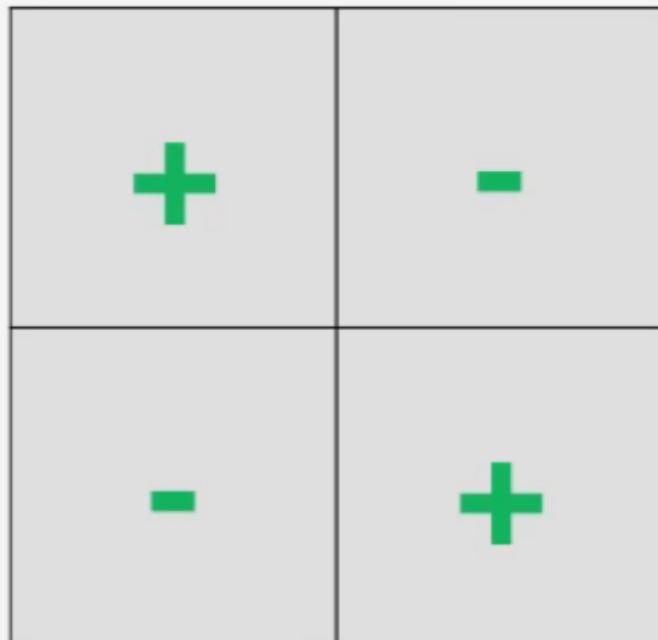


Image Classifier

If positive, “\”  
If negative, “/”

# What is convolution?

$$\begin{array}{|c|c|} \hline + & - \\ \hline 1 & -1 \\ \hline - & + \\ \hline -1 & 1 \\ \hline \end{array}$$

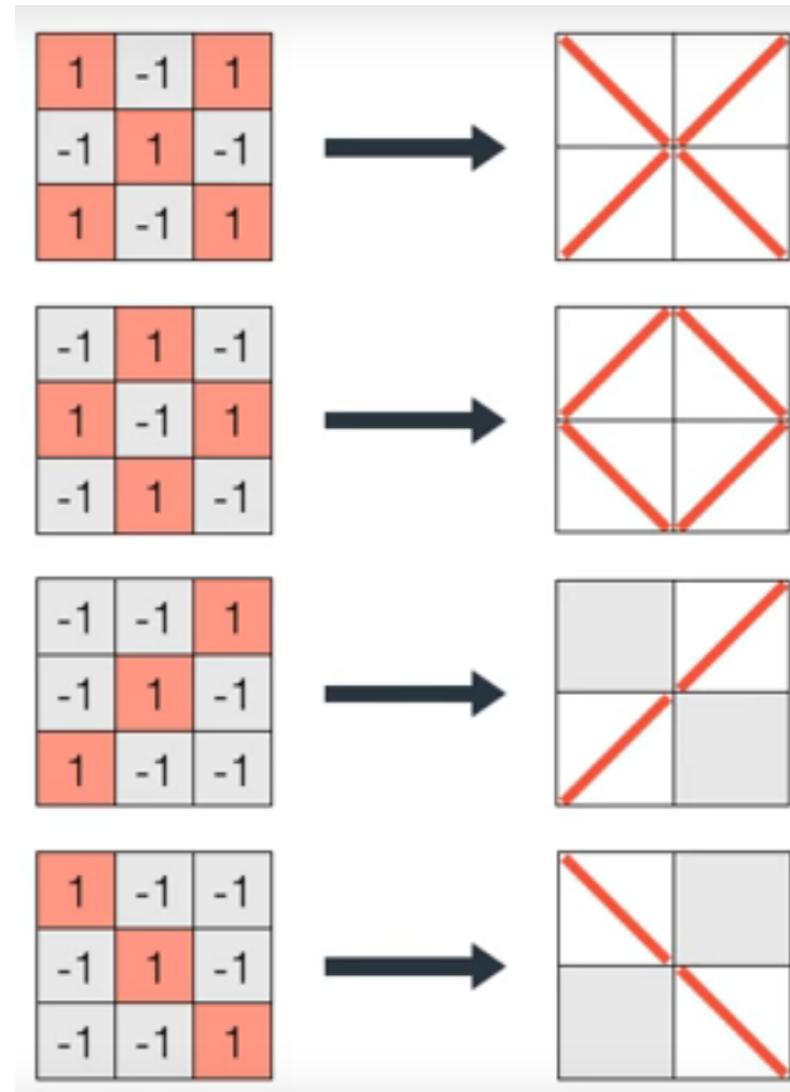
$$\begin{array}{cccc} + & - & - & + \\ \hline 1 & 1 & -1 & 1 \\ \hline +1 & -1 & +1 & +1 \\ \hline \end{array} = 2$$

$$\begin{array}{|c|c|} \hline + & - \\ \hline -1 & -1 \\ \hline - & + \\ \hline 1 & -1 \\ \hline \end{array}$$

$$\begin{array}{cccc} + & - & - & + \\ \hline -1 & -1 & 1 & -1 \\ \hline -1 & +1 & -1 & -1 \\ \hline \end{array} = -2$$

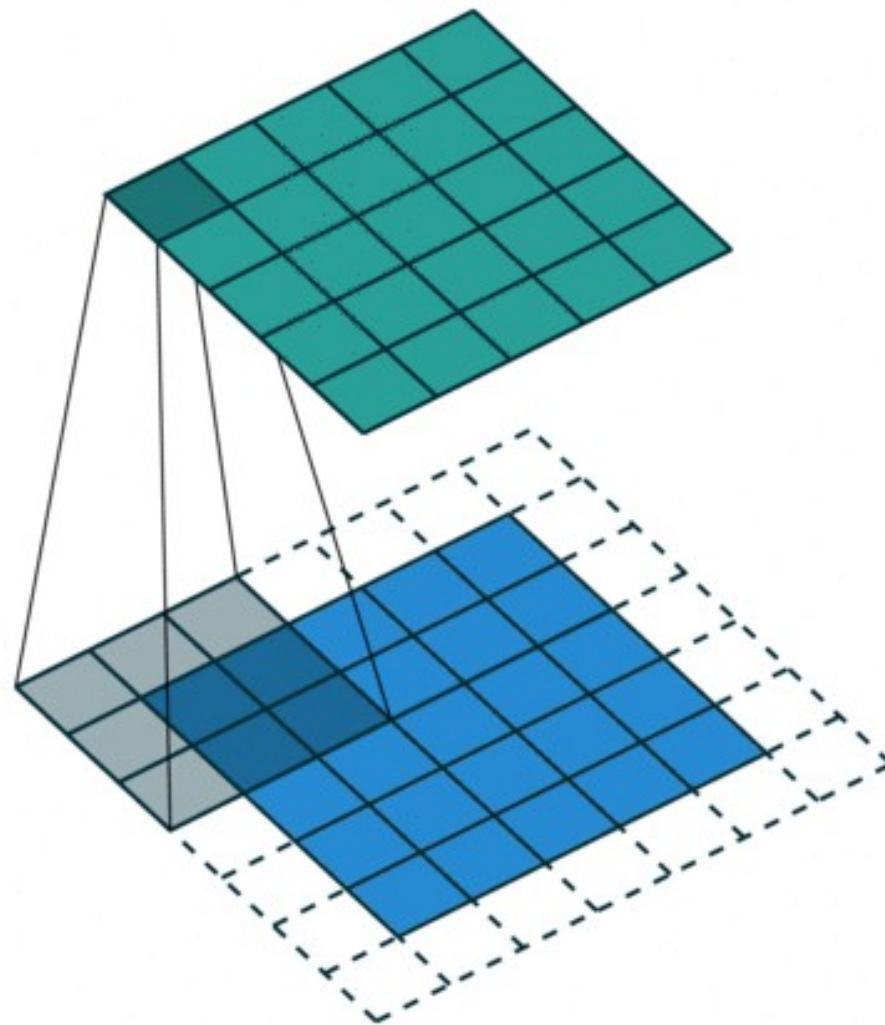
SOURCE: Luis Serrano, Udacity, <https://www.youtube.com/watch?v=2-Ol7ZB0MmU>

# What is convolution?



SOURCE: Luis Serrano, Udacity, <https://www.youtube.com/watch?v=2-Ol7ZB0MmU>

# What is convolution?



Feature Map



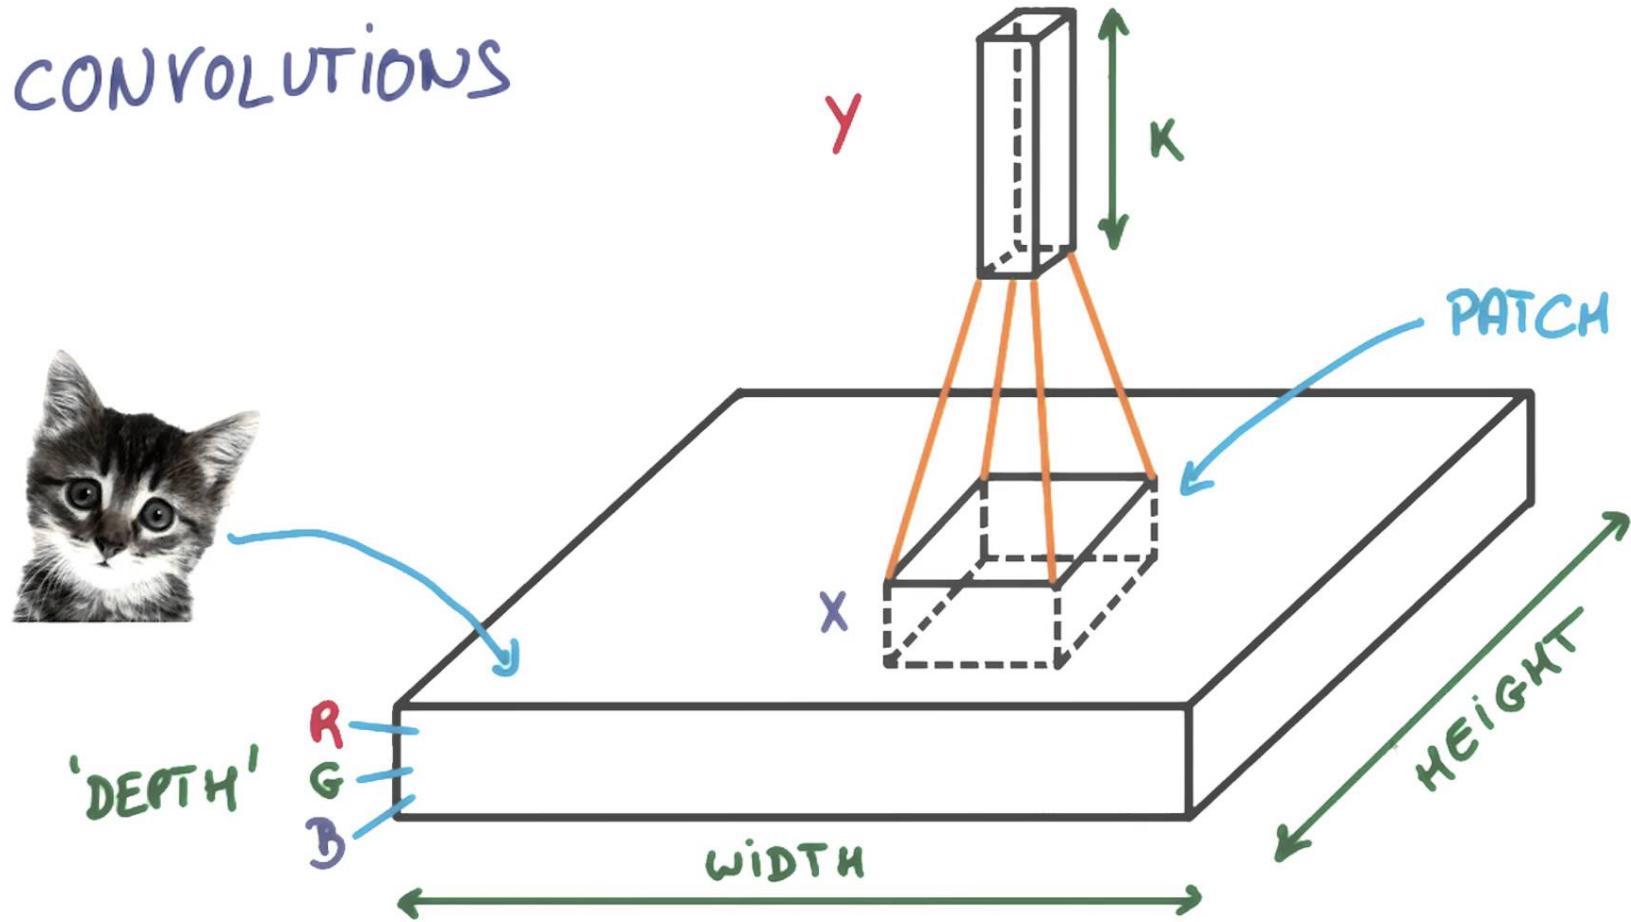
Kernel or Filter



Image

# What is convolution?

CONVOLUTIONS



Source: Vincent Vanhoucke, Udacity, <https://www.youtube.com/watch?v=jajksuQW4mc>

# What is convolution?

Image

3 <sub>0</sub>	3 <sub>1</sub>	2 <sub>2</sub>	1	0
0 <sub>2</sub>	0 <sub>2</sub>	1 <sub>0</sub>	3	1
3 <sub>0</sub>	1 <sub>1</sub>	2 <sub>2</sub>	2	3
2	0	0	2	2
2	0	0	0	1

Feature Map

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Source: Vincent Dumoulin and Francesco Visin, <https://arxiv.org/pdf/1603.07285.pdf>

# What is convolution?

Image

3	3 <sub>0</sub>	2 <sub>1</sub>	1 <sub>2</sub>	0
0	0 <sub>2</sub>	1 <sub>2</sub>	3 <sub>0</sub>	1
3	1 <sub>0</sub>	2 <sub>1</sub>	2 <sub>2</sub>	3
2	0	0	2	2
2	0	0	0	1

Feature  
Map

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Source: Vincent Dumoulin and Francesco Visin, <https://arxiv.org/pdf/1603.07285.pdf>

# What is convolution?

Image

3	3	2 <sub>0</sub>	1 <sub>1</sub>	0 <sub>2</sub>
0	0	1 <sub>2</sub>	3 <sub>2</sub>	1 <sub>0</sub>
3	1	2 <sub>0</sub>	2 <sub>1</sub>	3 <sub>2</sub>
2	0	0	2	2
2	0	0	0	1

Feature Map

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Source: Vincent Dumoulin and Francesco Visin, <https://arxiv.org/pdf/1603.07285.pdf>

# What is convolution?

Image

3	3	2	1	0
0 <sub>0</sub>	0 <sub>1</sub>	1 <sub>2</sub>	3	1
3 <sub>2</sub>	1 <sub>2</sub>	2 <sub>0</sub>	2	3
2 <sub>0</sub>	0 <sub>1</sub>	0 <sub>2</sub>	2	2
2	0	0	0	1

Feature  
Map

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Source: Vincent Dumoulin and Francesco Visin, <https://arxiv.org/pdf/1603.07285.pdf>

# What is convolution?

Image

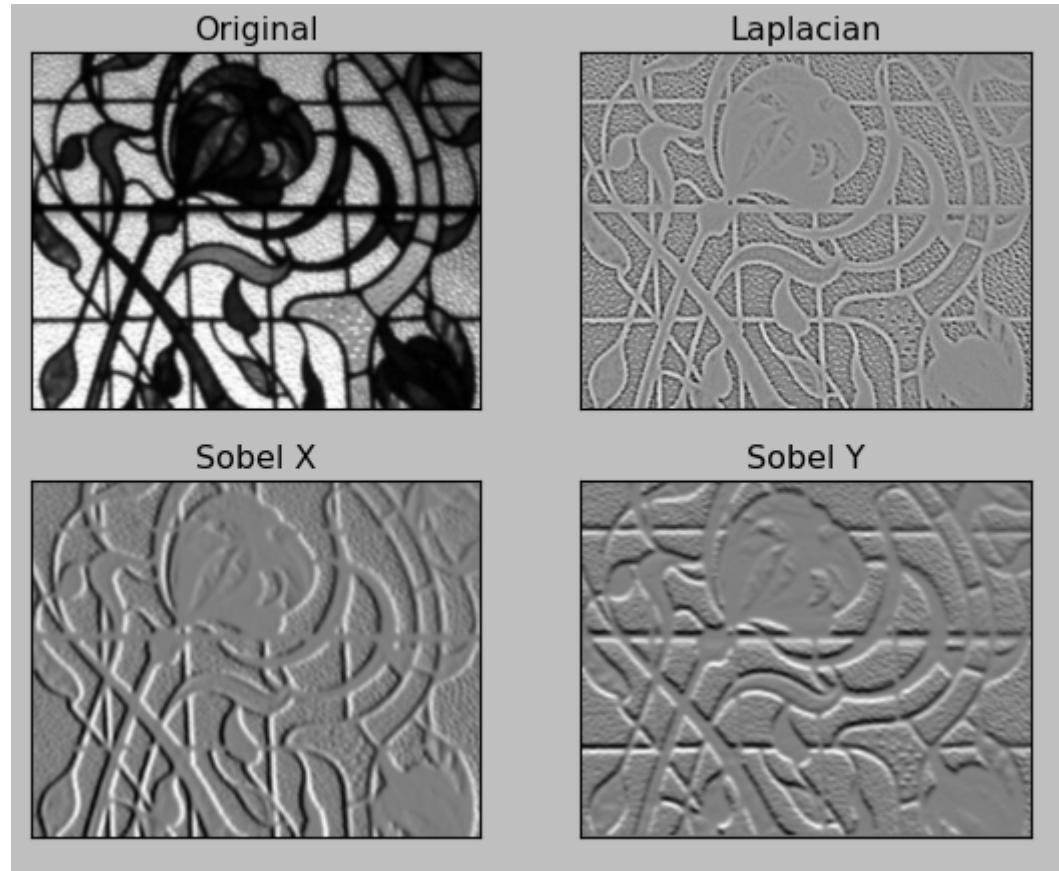
3	3	2	1	0
0	0 <sub>0</sub>	1 <sub>1</sub>	3 <sub>2</sub>	1
3	1 <sub>2</sub>	2 <sub>2</sub>	2 <sub>0</sub>	3
2	0 <sub>0</sub>	0 <sub>1</sub>	2 <sub>2</sub>	2
2	0	0	0	1

Feature  
Map

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Source: Vincent Dumoulin and Francesco Visin, <https://arxiv.org/pdf/1603.07285.pdf>

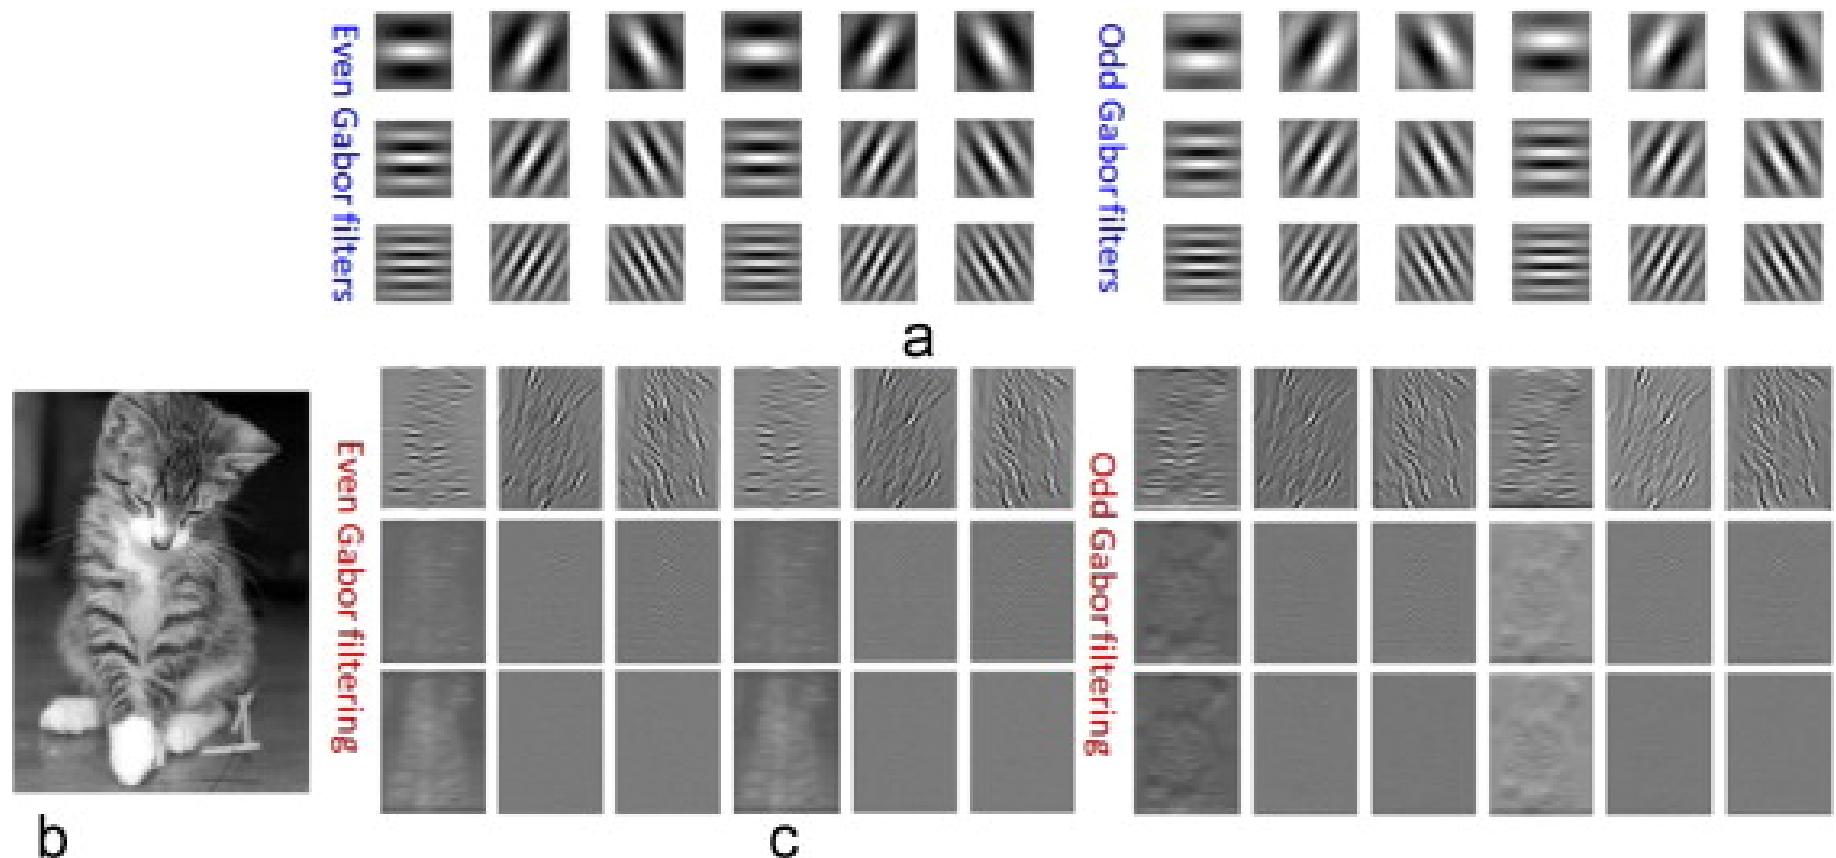
Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur $3 \times 3$ (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	
Gaussian blur $5 \times 5$ (approximation)	$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	



Watch the video: <http://setosa.io/ev/image-kernels/>

# What is convolution?

## Gabor Filter



Source: <http://www.sciencedirect.com/science/article/pii/S0031320315001570>

# What is convolution?

a convolution matrix

22	15	1	3	60
42	5	38	39	7
28	9	4	66	79
0	82	45	12	17
99	14	72	51	3

$\times$

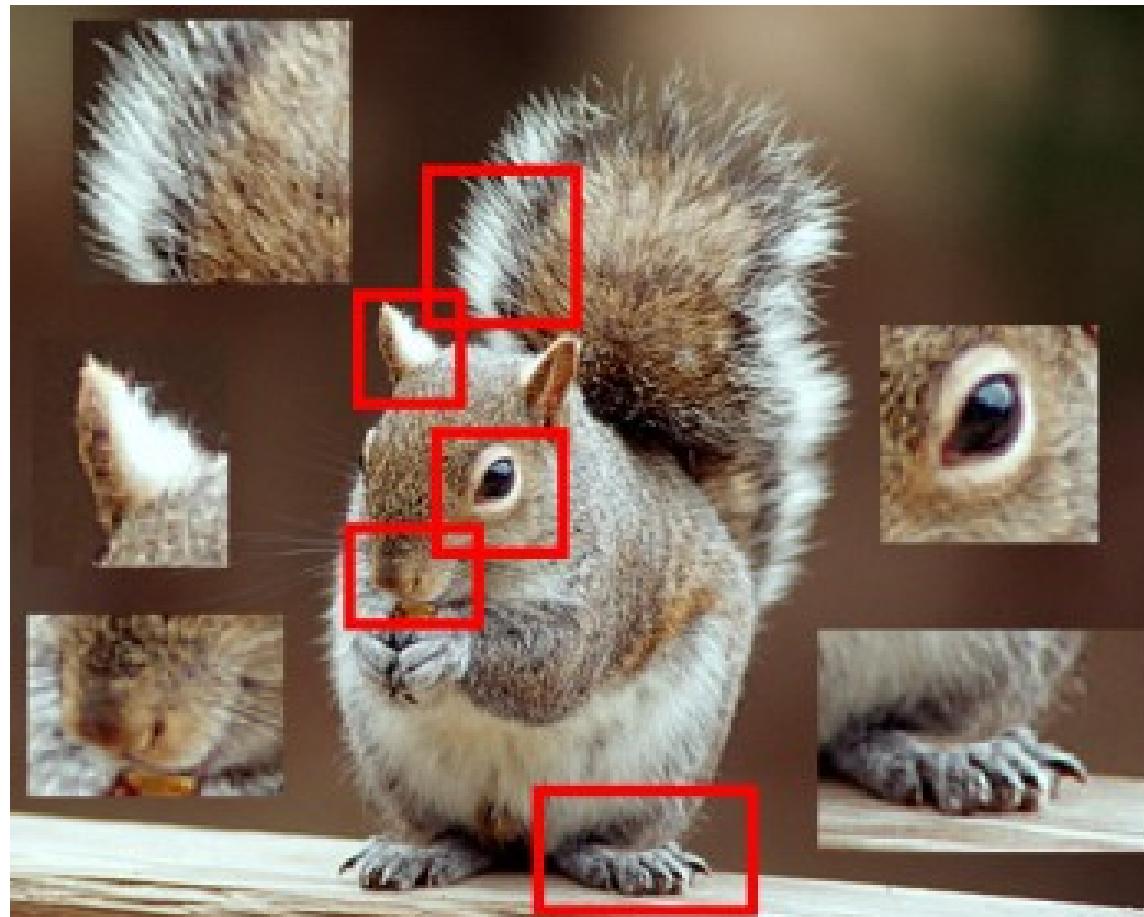
0 0 0 0 0  
0 0 0 1 0  
0 0 0 0 0  
0 0 0 0 0  
0 0 0 0 0

=

	1	3	60	
	38	39	7	
	4	66	79	

“Take the top right corner of the image?”

# What is convolution?



Convolution creates FEATURE DETECTORS.

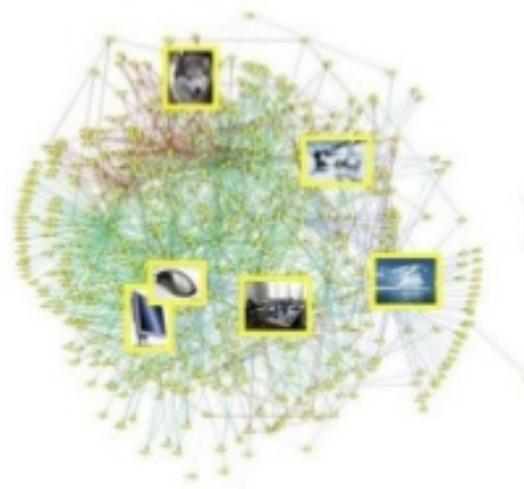
# ImageNet

ImageNet

2010-2017



Prof. Fei-Fei Li  
Stanford University  
Creator of ImageNet



IMAGENET

<https://qz.com/1034972/the-data-that-changed-the-direction-of-ai-research-and-possibly-the-world/>

# PASCAL VOC

Aeroplanes  
Bicycles  
Birds  
Boats  
Bottles  
Buses  
Cars  
Cats  
Chairs  
Cows  
Dining tables  
Dogs  
Horses  
Motorbikes  
People  
Potted plants  
Sheep  
Sofas  
Trains  
TV/Monitors

21,738

# ImageNet

High level category	# synset (subcategories)	Avg # images per synset	Total # images
amphibian	94	591	56K
animal	3822	732	2799K
appliance	51	1164	59K
bird	856	949	812K
covering	946	819	774K
device	2385	675	1610K
fabric	262	690	181K
fish	566	494	280K
flower	462	735	339K
food	1495	670	1001K
fruit	309	607	188K
fungus	303	453	137K
furniture	187	1043	195K
geological formation	151	838	127K
invertebrate	728	573	417K
mammal	1138	821	934K
musical instrument	157	891	140K
plant	1666	600	999K
reptile	268	707	190K
sport	166	1207	200K
structure	1239	763	946K
tool	316	551	174K
tree	993	568	564K
utensil	86	912	78K
vegetable	176	764	135K
vehicle	481	778	374K
person	2035	468	952K

14,197,122

# ImageNet

## English foxhound

An English breed slightly larger than the American foxhounds originally used to hunt in packs

454  
pictures

37.57%  
Popularity  
Percentile



- survivor (0)
- range animal (0)
- creepy-crawly (0)
- domestic animal, domesticated animal (213)
  - domestic cat, house cat, *Felis domesticus*, *Felis catus* (18)
  - dog, domestic dog, *Canis familiaris* (189)
    - pooch, doggie, doggy, barker, bow-wow (0)
    - hunting dog (101)
      - sporting dog, gun dog (28)
      - dachshund, dachsie, badger dog (1)
      - terrier (37)
      - courser (0)
      - hound, hound dog (29)
        - Plott hound (0)
        - wolfhound (2)
        - Scottish deerhound, deerhound (0)
        - coonhound (2)
        - foxhound (3)
          - Walker hound, Walker foxhound (0)
          - American foxhound (0)
          - English foxhound (0)
        - Weimaraner (0)
        - otterhound, otter hound (0)
        - bloodhound, sleuthhound (0)
        - Norwegian elkhound, elkhound (0)
        - Saluki, gazelle hound (0)
        - Afghan hound, Afghan (0)
        - staghound (0)
        - greyhound (2)
        - beagle (0)
        - harrier (0)
        - basset, basset hound (0)
        - bluetick (0)
        - redbone (0)

Treemap VisualizationImages of the SynsetDownloads







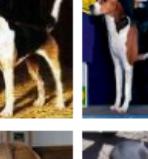




























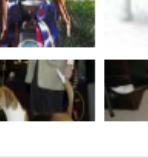
















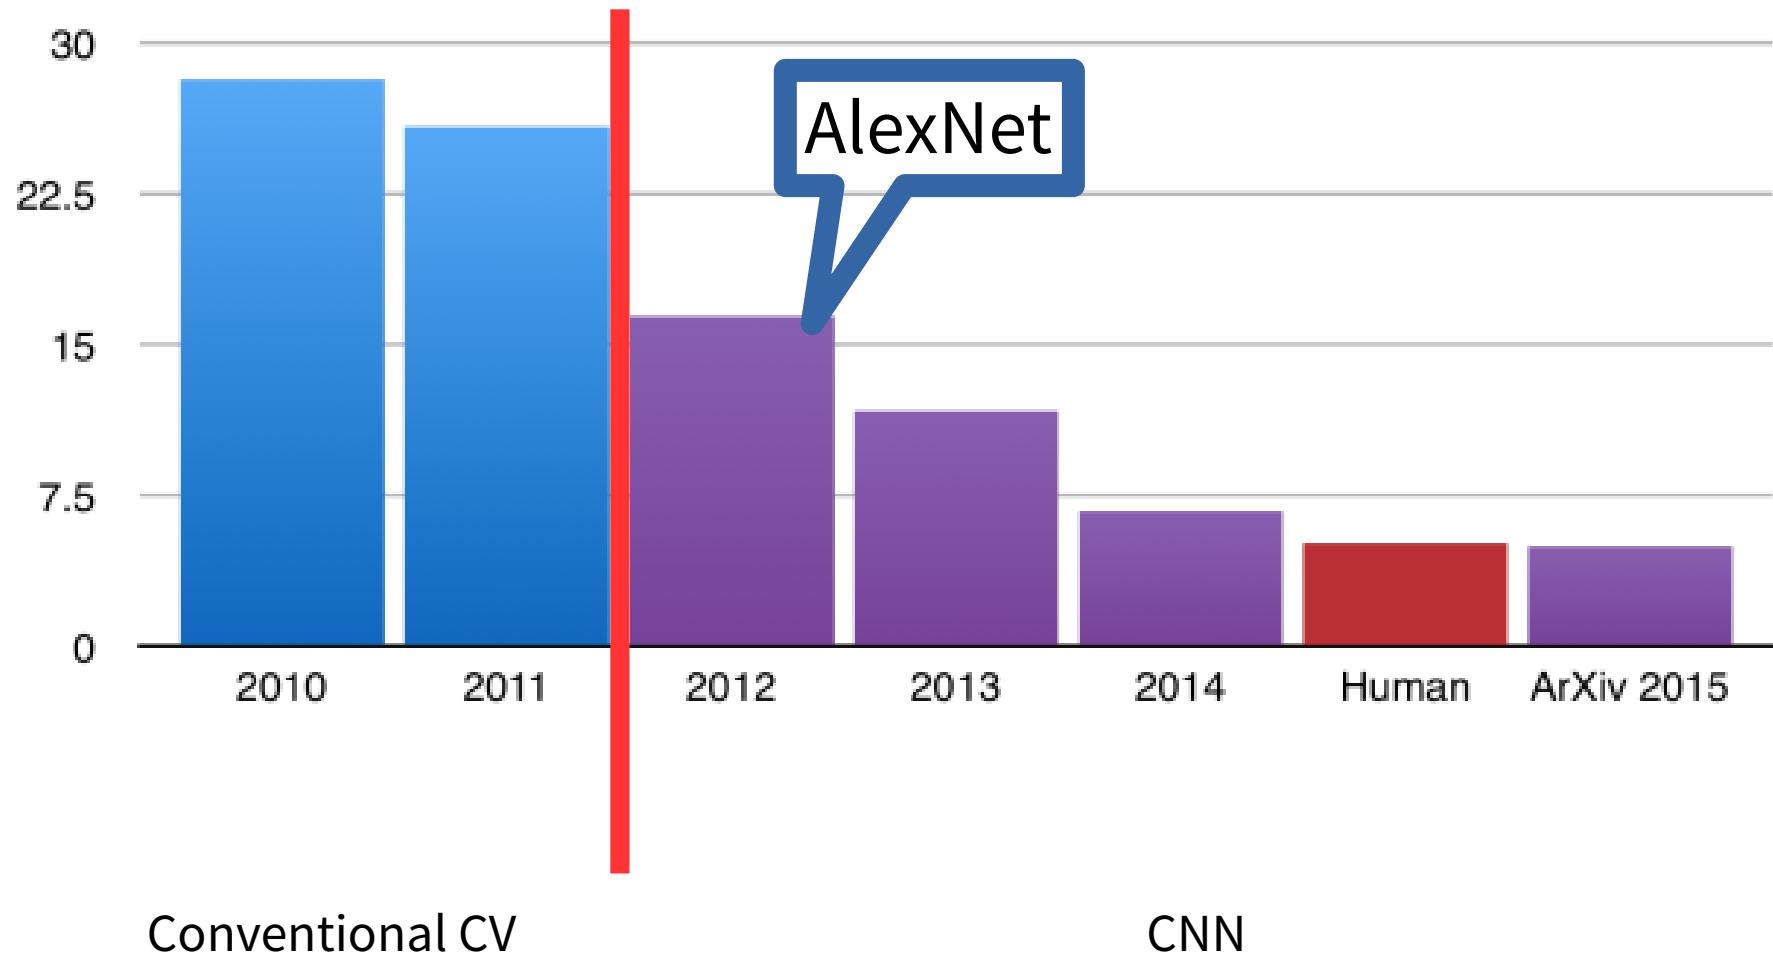
# ImageNet



Source: <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks>

# ImageNet

ILSVRC top-5 error on ImageNet

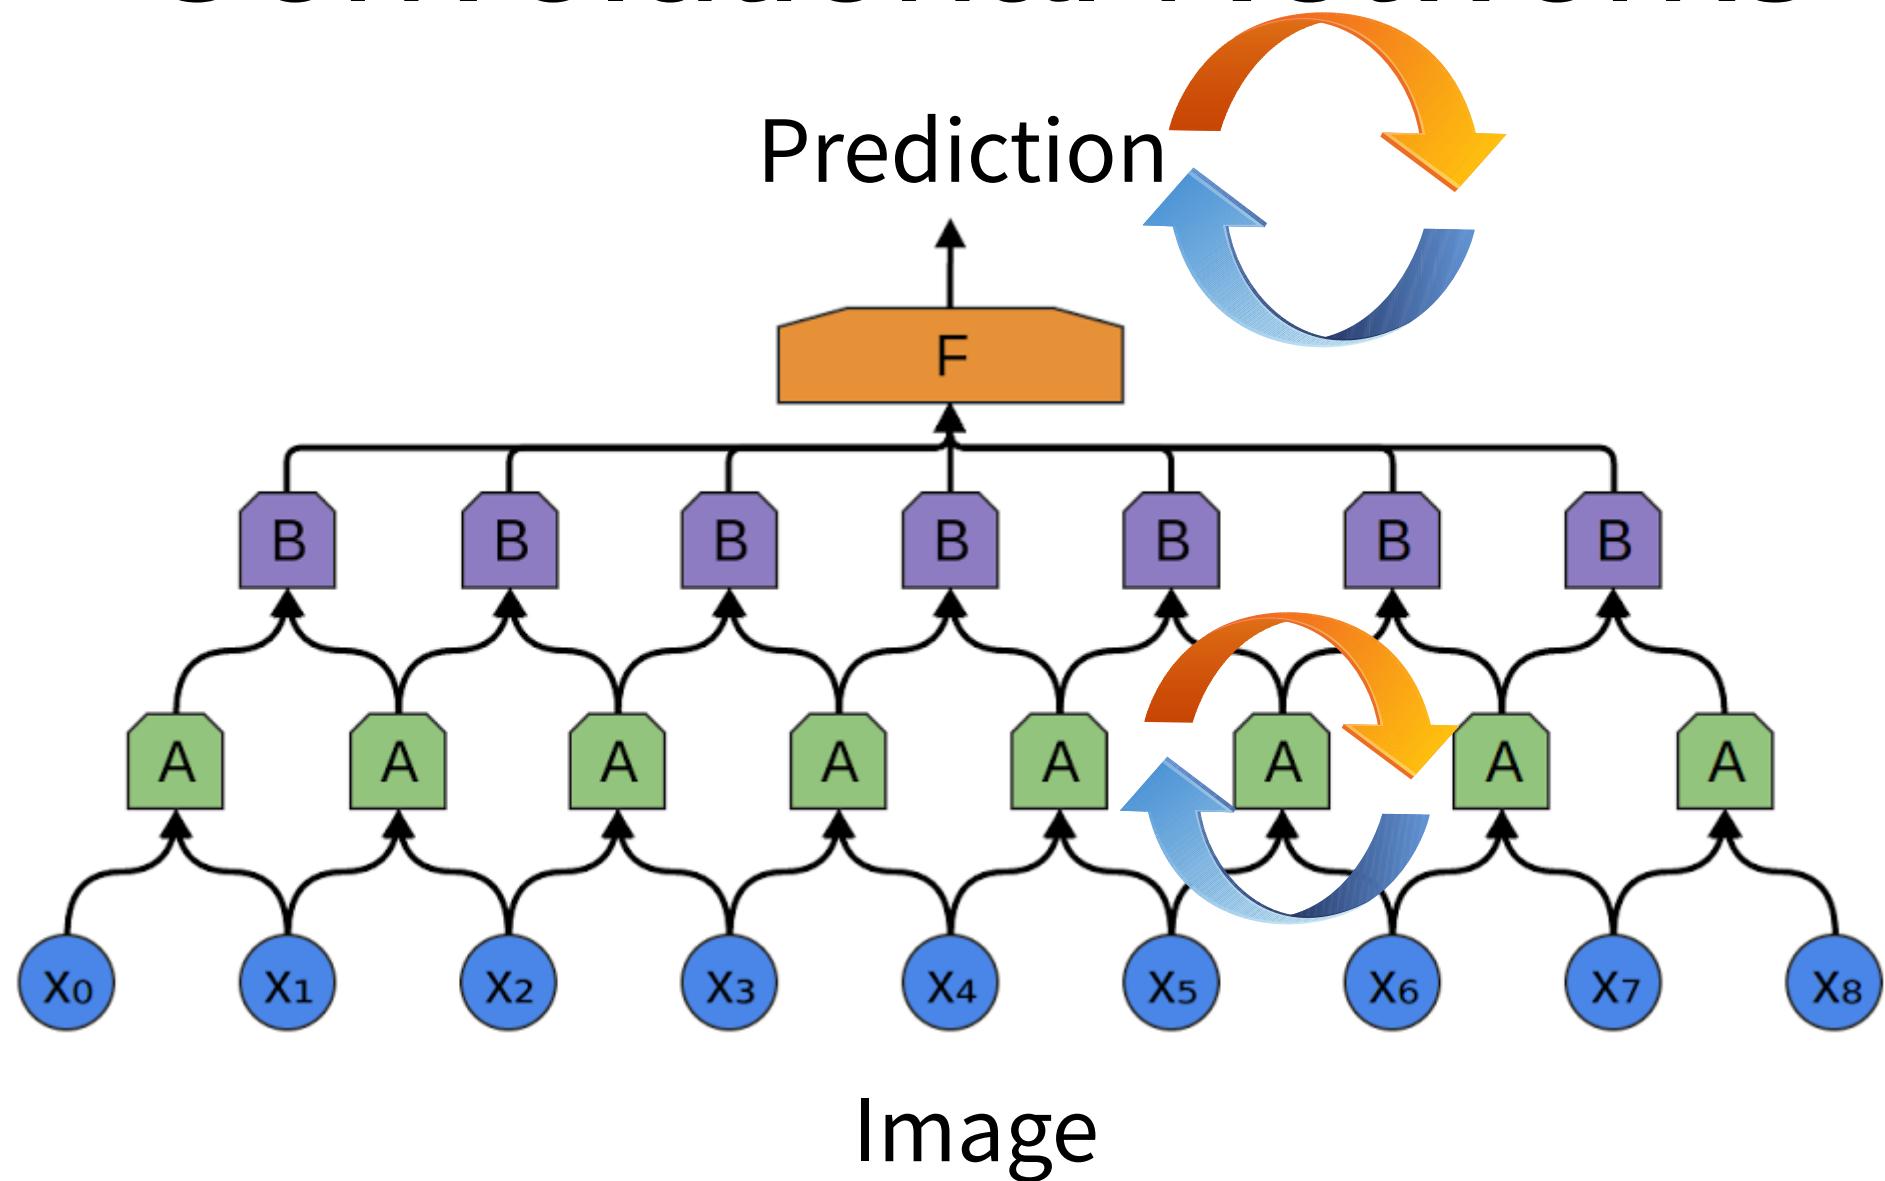


# Convolutional Networks

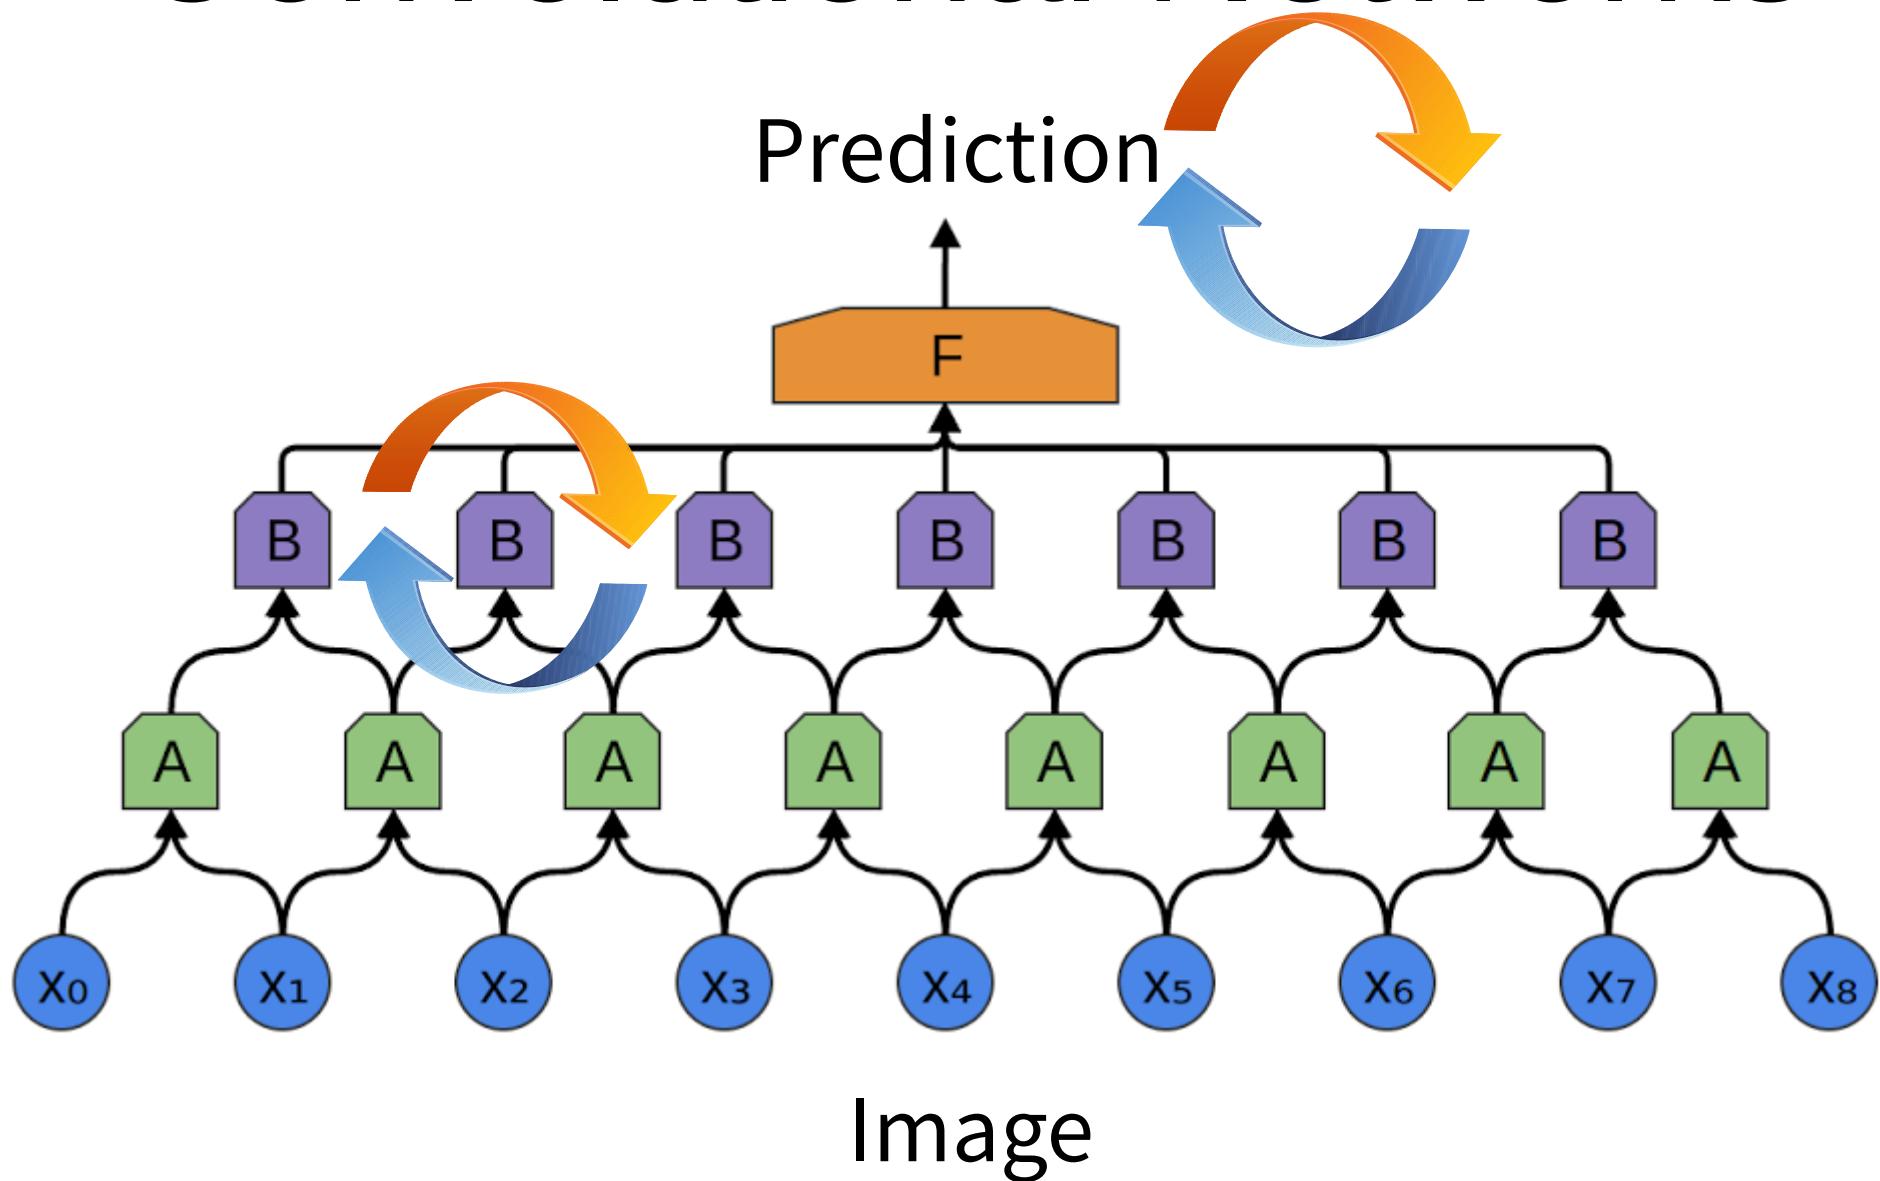
## AlexNet

- Convolutional Network
- Dropout
- Rectified Linear Unit (ReLU)
- Backpropagation
- Stochastic Gradient Descent

# Convolutional Networks

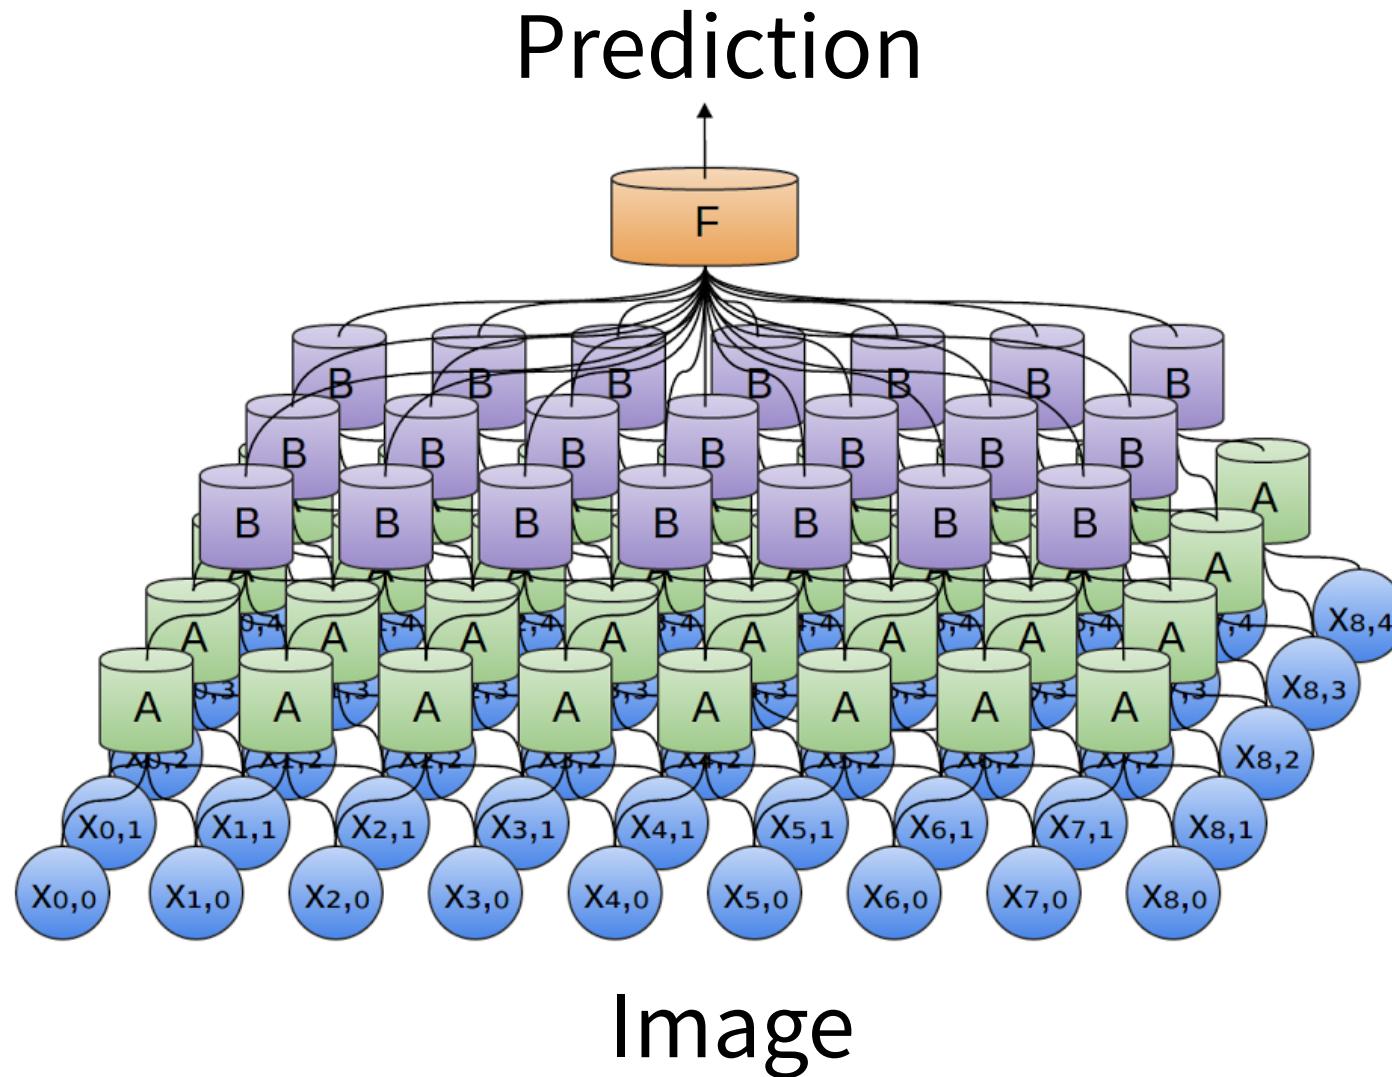


# Convolutional Networks

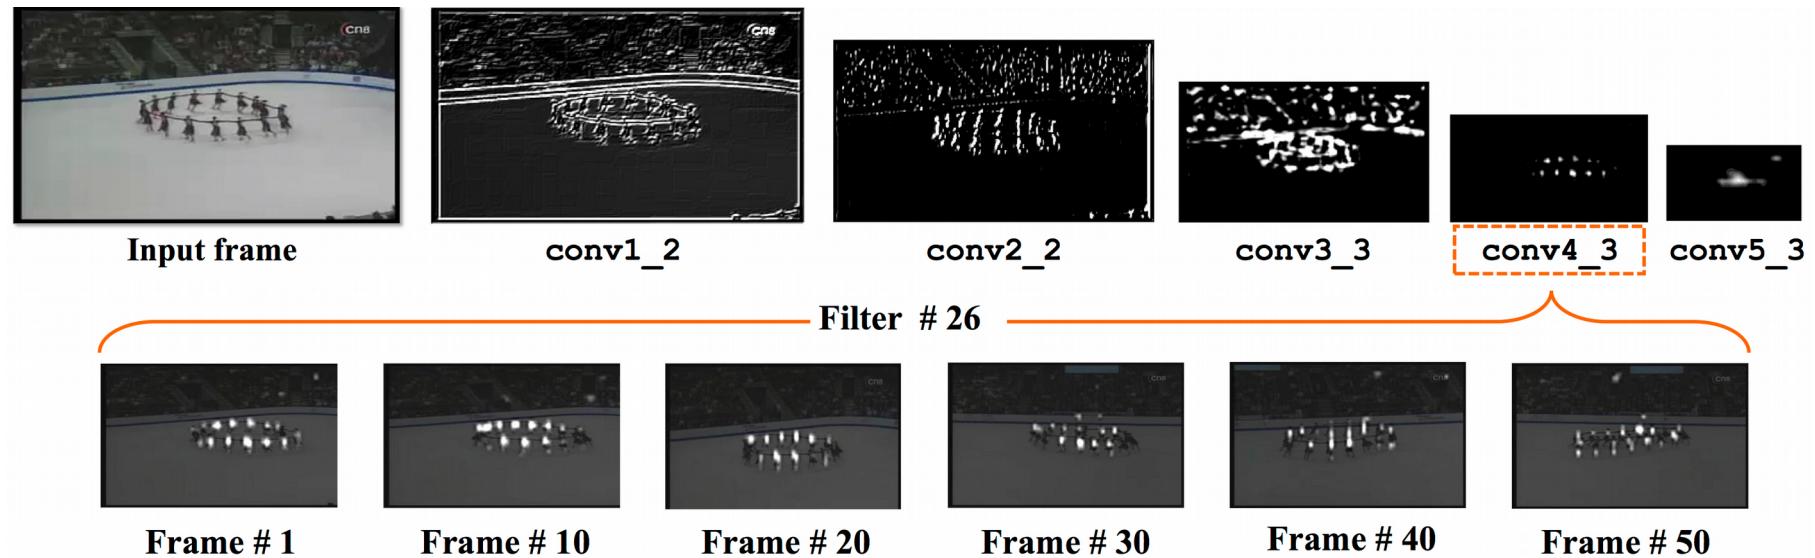


Backpropagation: <http://colah.github.io/posts/2015-08-Backprop/>

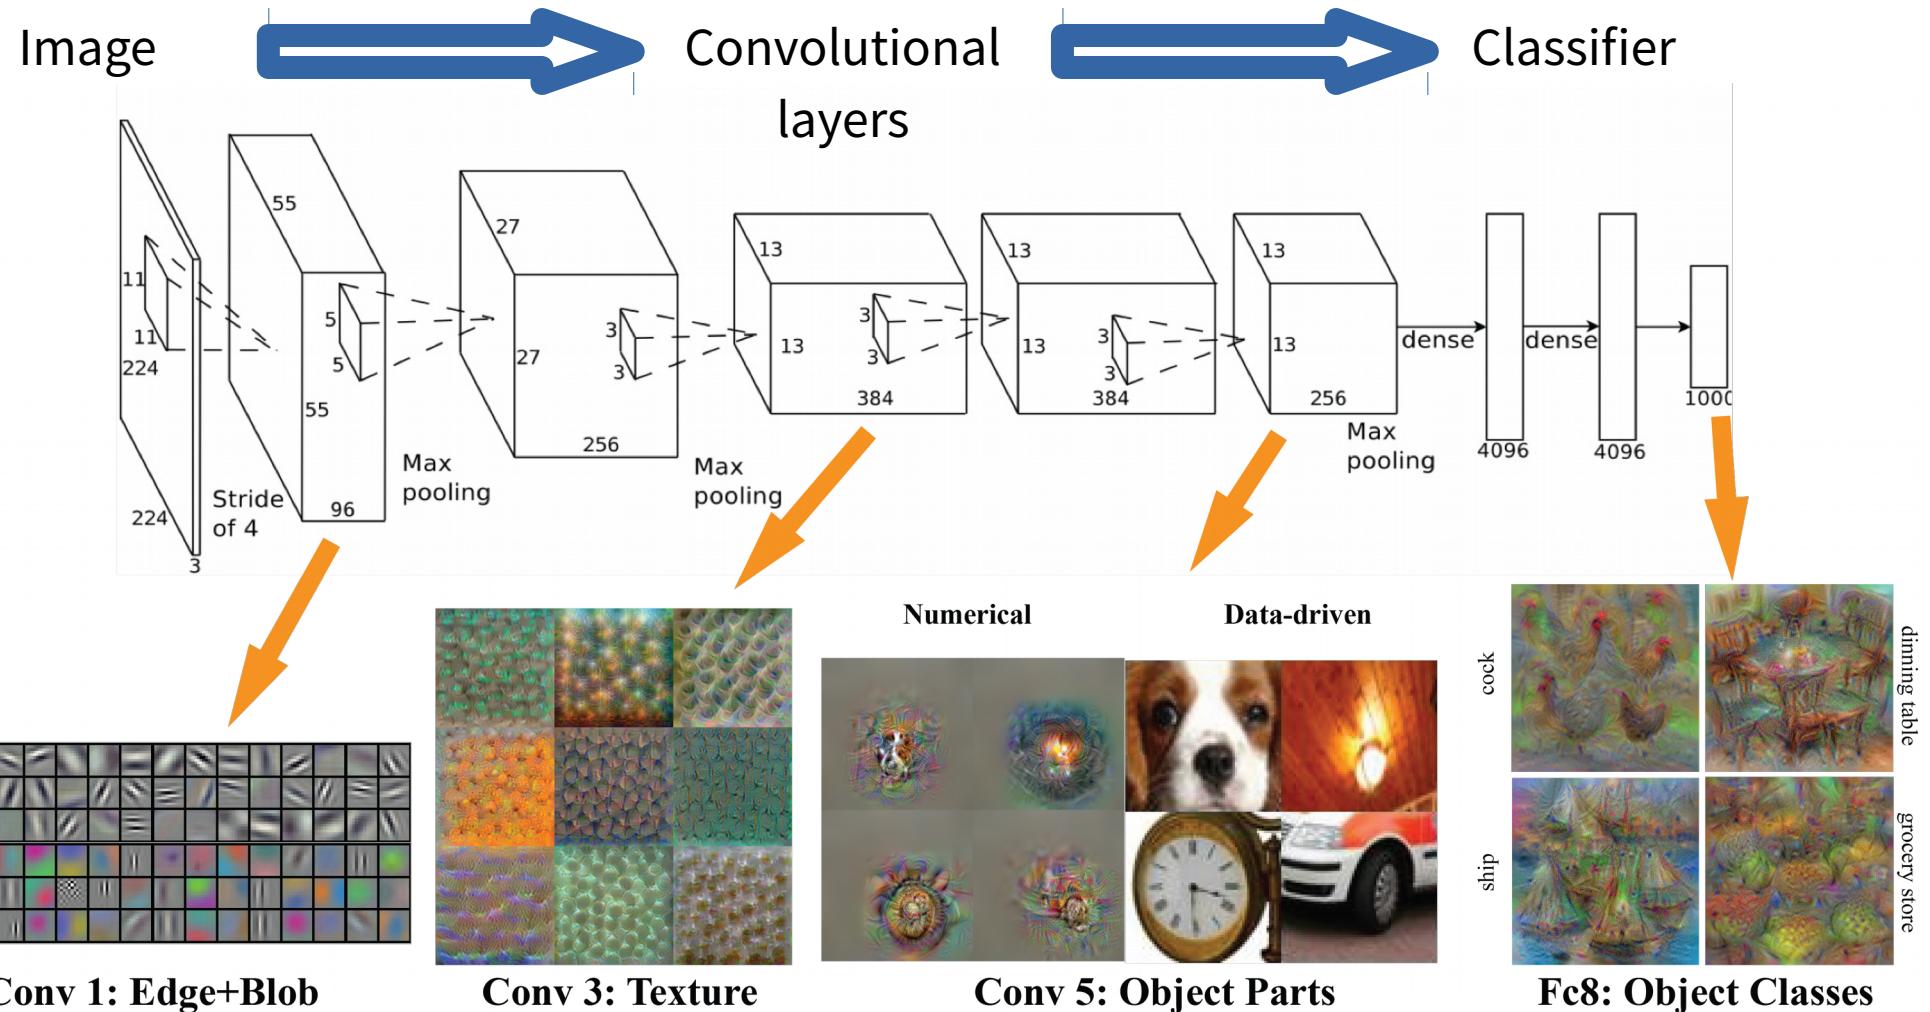
# Convolutional Networks



# Convolutional Networks



# Convolutional Networks



Source: [http://vision03.csail.mit.edu/cnn\\_art/data/single\\_layer.png](http://vision03.csail.mit.edu/cnn_art/data/single_layer.png)

# Convolutional Networks

## Deep Visualization Toolbox

[yosinski.com/deepvis](http://yosinski.com/deepvis)

#deepvis



Jason Yosinski



Jeff Clune



Anh Nguyen



Thomas Fuchs



Hod Lipson



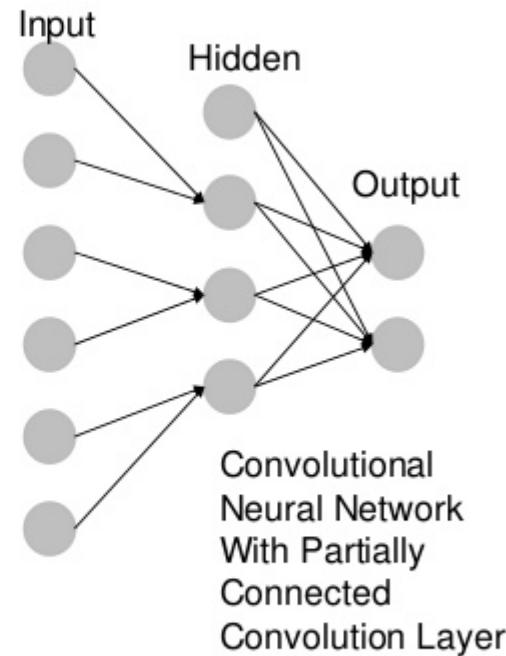
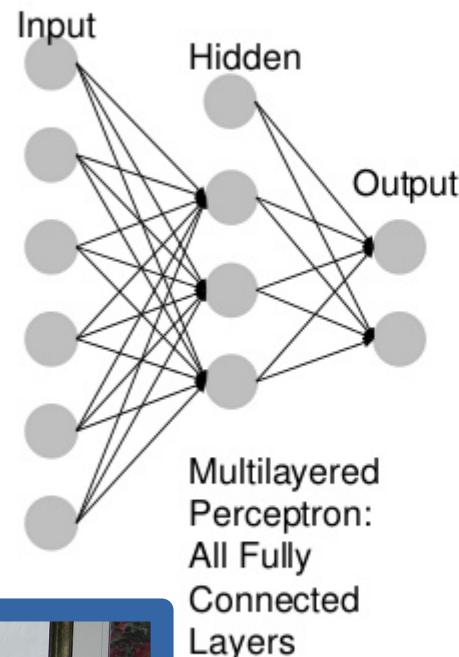
Cornell University



**Jet Propulsion Laboratory**  
California Institute of Technology

# Convolutional Networks

## MLP VS ConvNet



# Convolutional Networks

MLP  
98.2%

CNN  
99.3%

Layer (type)	Output Shape	Param #
<hr/>		
dense_1 (Dense)	(None, 784)	615440
dense_2 (Dense)	(None, 10)	7850
<hr/>		
Total params: 623,290		
Trainable params: 623,290		
Non-trainable params: 0		

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 30, 24, 24)	780
max_pooling2d_1 (MaxPooling2D)	(None, 30, 12, 12)	0
conv2d_2 (Conv2D)	(None, 15, 10, 10)	4065
max_pooling2d_2 (MaxPooling2D)	(None, 15, 5, 5)	0
dropout_1 (Dropout)	(None, 15, 5, 5)	0
flatten_1 (Flatten)	(None, 375)	0
dense_1 (Dense)	(None, 128)	48128
dense_2 (Dense)	(None, 50)	6450
dense_3 (Dense)	(None, 10)	510
<hr/>		
Total params: 59,933		
Trainable params: 59,933		
Non-trainable params: 0		

# Code Example in Keras

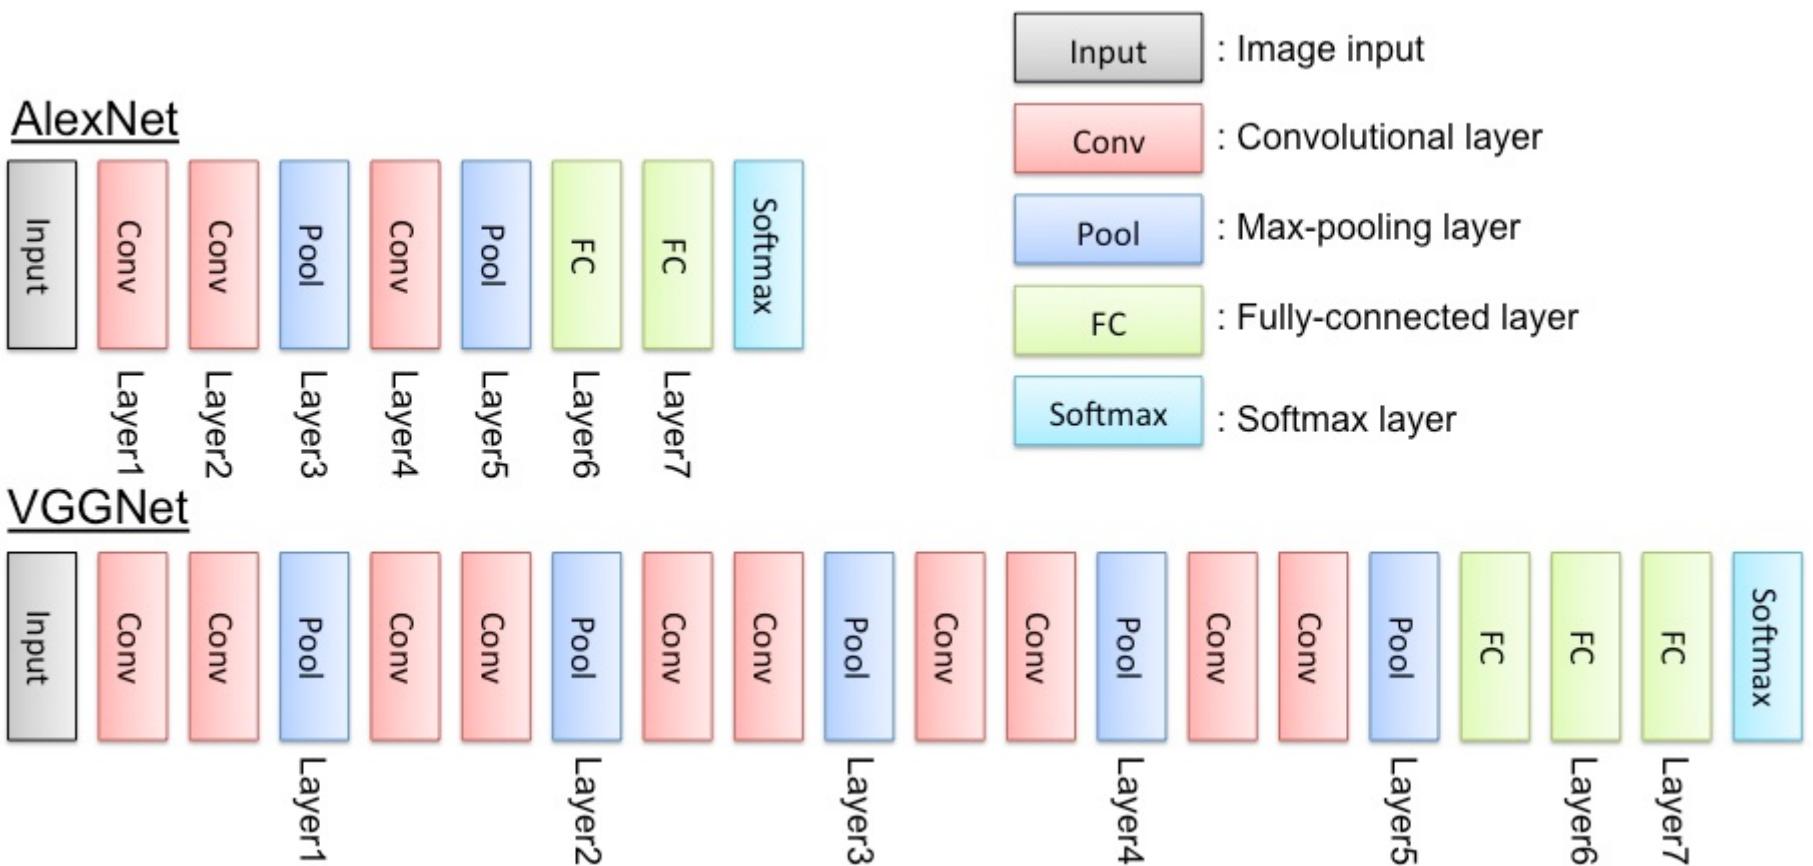
```
# define the convolutional model
def conv_model():

    model = Sequential()

    # input_shape tells Keras/Tensorflow to
    #     expect a tensor input of 1 x 28 x 28
    model.add(Conv2D(30, (5, 5), input_shape=(1, 28, 28), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Conv2D(15, (3, 3), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.2))
    model.add(Flatten())
    model.add(Dense(128, activation='relu'))
    model.add(Dense(50, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))

    # Compile model
    model.compile(loss='categorical_crossentropy',
                  optimizer='adam', metrics=['accuracy'])
    return model
```

# AlexNet v. VGG

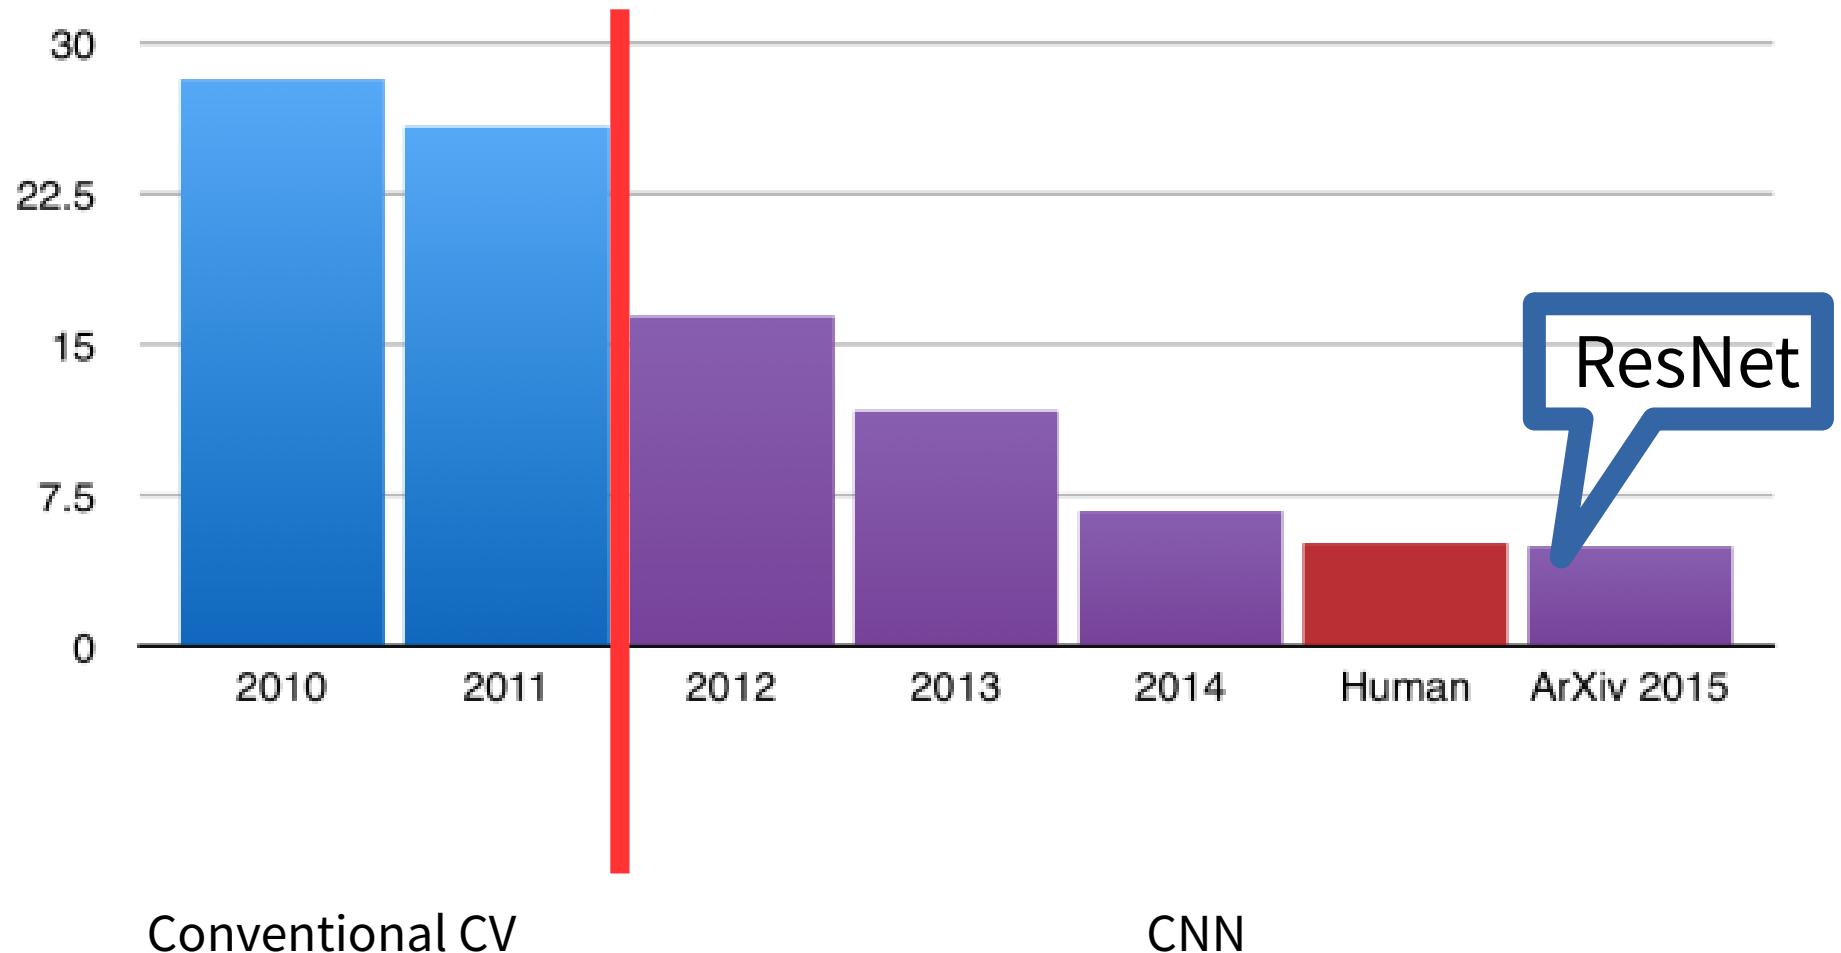


# Convolutional Networks

```
# Set up the model layers
vgg_layers = []
# set up 3x3 conv stacks with different number of filters
vgg_layers.append(Conv((3, 3, 64), **conv_params))
vgg_layers.append(Conv((3, 3, 64), **conv_params))
vgg_layers.append(Pooling(2, strides=2))
vgg_layers.append(Conv((3, 3, 128), **conv_params))
vgg_layers.append(Conv((3, 3, 128), **conv_params))
vgg_layers.append(Pooling(2, strides=2))
vgg_layers.append(Conv((3, 3, 256), **conv_params))
vgg_layers.append(Conv((3, 3, 256), **conv_params))
vgg_layers.append(Pooling(2, strides=2))
vgg_layers.append(Conv((3, 3, 512), **conv_params))
vgg_layers.append(Conv((3, 3, 512), **conv_params))
vgg_layers.append(Pooling(2, strides=2))
vgg_layers.append(Conv((3, 3, 512), **conv_params))
vgg_layers.append(Conv((3, 3, 512), **conv_params))
vgg_layers.append(Pooling(2, strides=2))
vgg_layers.append(Affine(nout=4096, init=GlorotUniform(), bias=Constant(0), activation=relu))
vgg_layers.append(Dropout(keep=0.5))
vgg_layers.append(Affine(nout=4096, init=GlorotUniform(), bias=Constant(0), activation=relu))
vgg_layers.append(Dropout(keep=0.5))
vgg_layers.append(Affine(nout=1, init=GlorotUniform(), bias=Constant(0), activation=Logistic(), name="class_layer"))
```

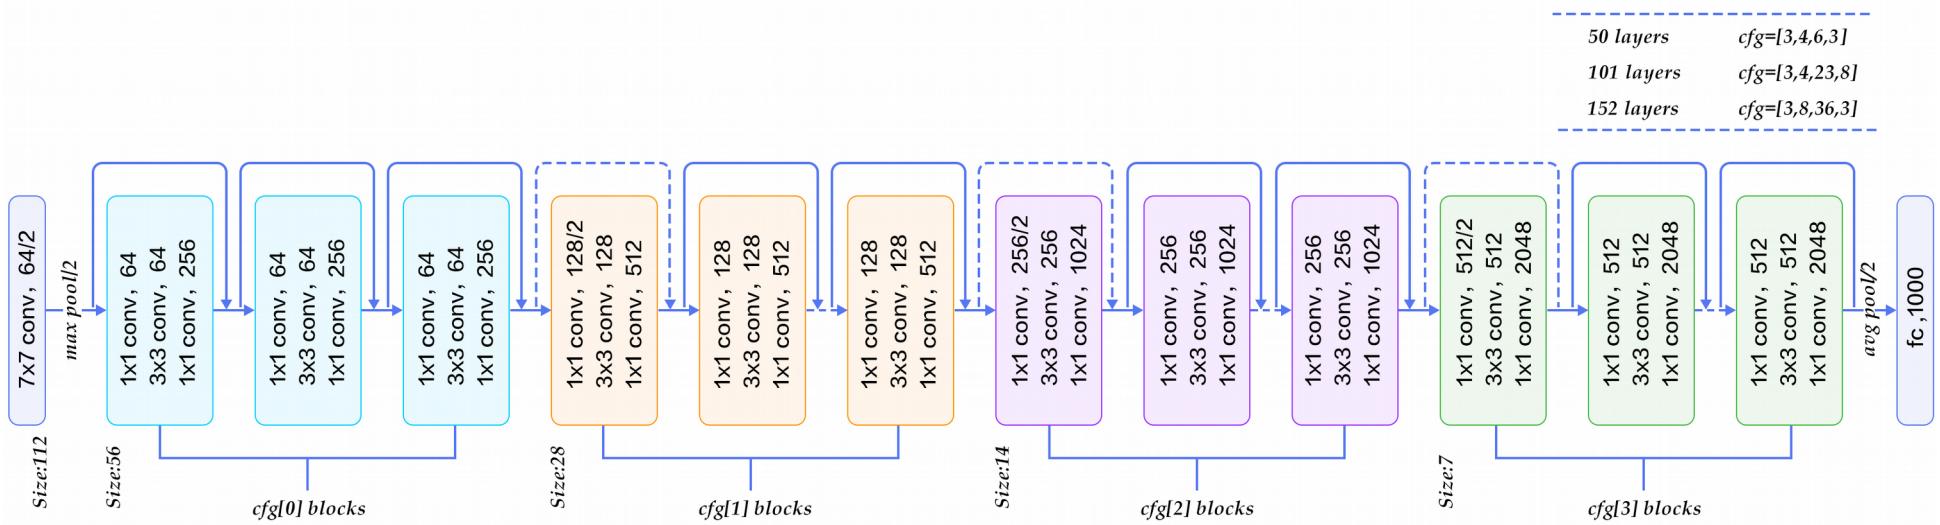
# ImageNet

# ILSVRC top-5 error on ImageNet



# ResNet

## *The Residual Network*

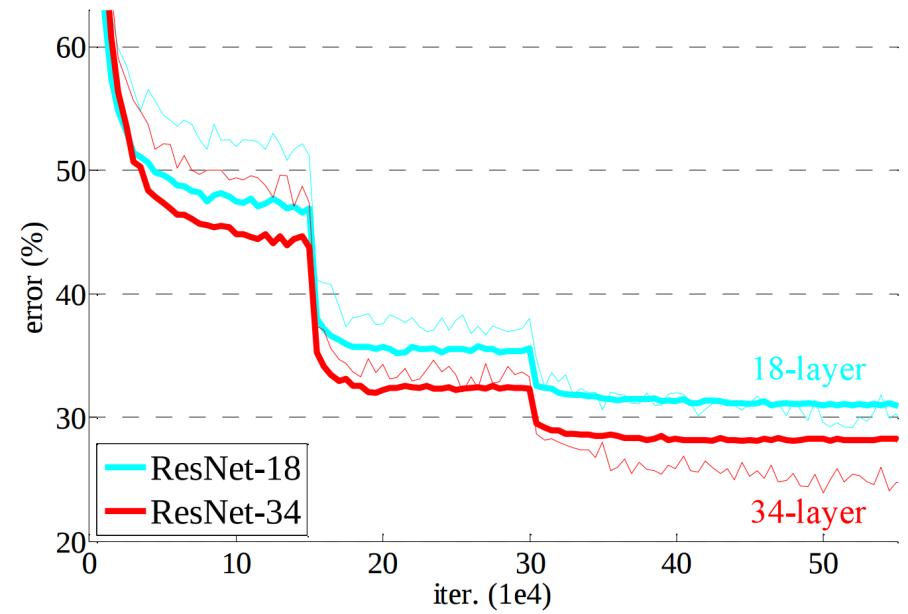
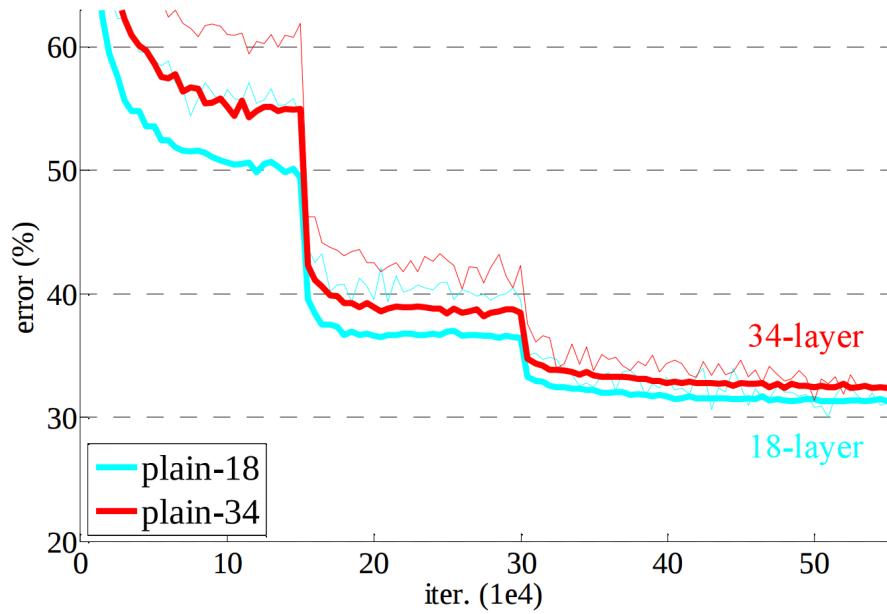


Kaiming He, Microsoft Research Asia (now at Facebook with Yann LeCun)  
<https://arxiv.org/abs/1512.03385>

[https://github.com/mas-dse-greina/neon/blob/master/luna16/LUNA16\\_resnet\\_HDF5.py](https://github.com/mas-dse-greina/neon/blob/master/luna16/LUNA16_resnet_HDF5.py)

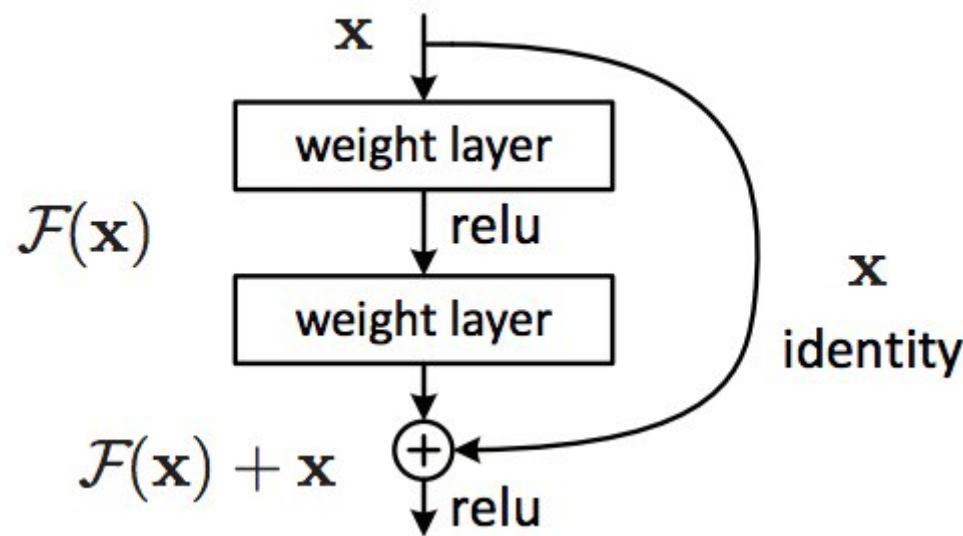
# ResNet

## *The Residual Network*



# ResNet

## *The Residual Network*



*Let the network “decide”  
how deep it needs to be.*

*For a **much** better lecture...*

<https://www.youtube.com/watch?v=LxfUGhugiQ&list=PLkt2uSq6rBVctENoVBg1TpCC7OQi31AlC&index=7>

Andrej Karpathy's blog:  
<http://karpathy.github.io/>